

Amazon Kinesis Data Streams



Fernando Medina Corey

CLOUD ARCHITECT

@fmc_sea www.fernandomc.com



Overview



Core Kinesis Data Streams concepts

- Producers, shards, consumers, and records

Producers

- Kinesis Agent, Producer Library, and AWS SDKs

Consumers

- Kinesis Client Libraries, AWS Services, and AWS SDKs

Demo - Analyzing Kinesis Data Streams



Kinesis Data Streams Basics




Producers



Kinesis Agent



Producer Library



AWS SDKs




Kinesis Data Streams Concepts



Data Stream



**Kinesis Data
Stream Shards**




Records



Consumers



Client Libraries



AWS SDKs



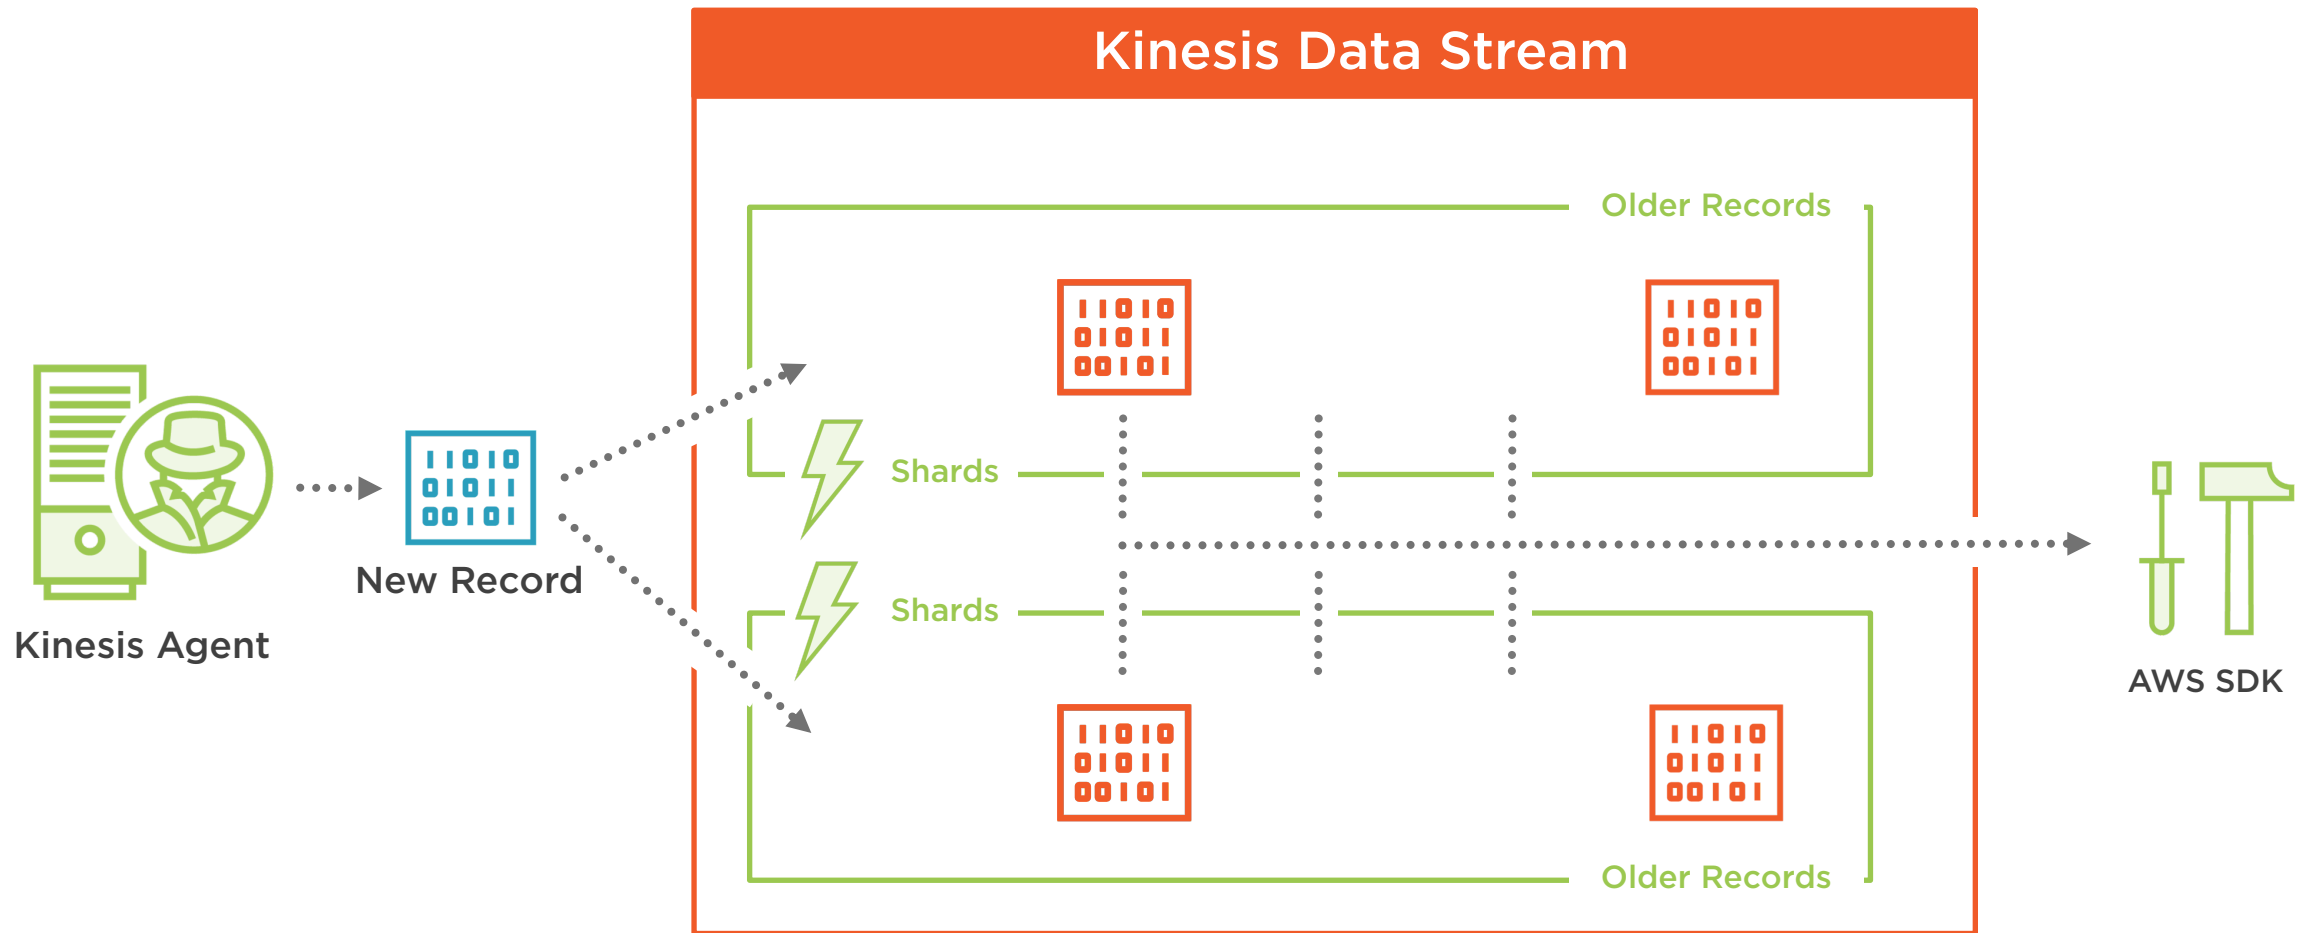
AWS Services



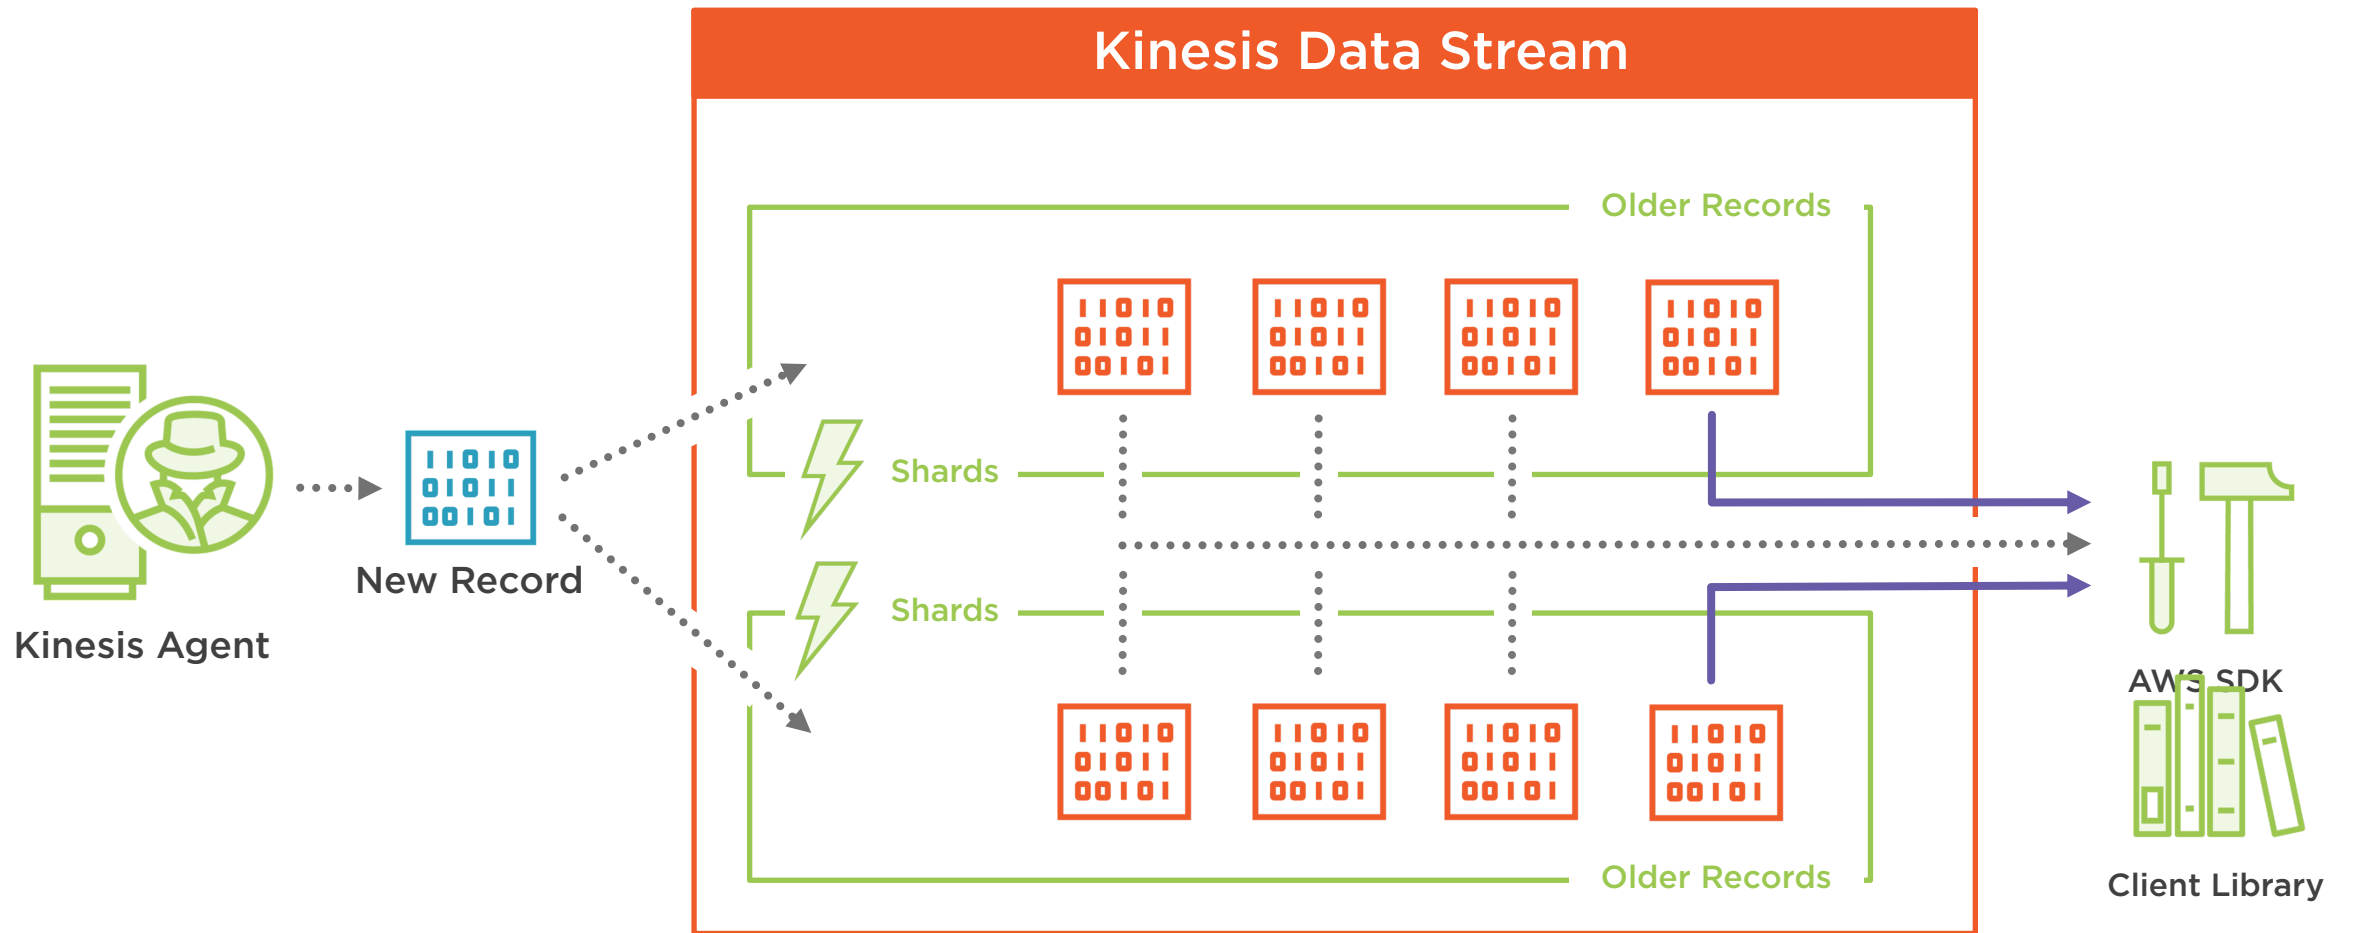
Kinesis Data Streams Shards and Records



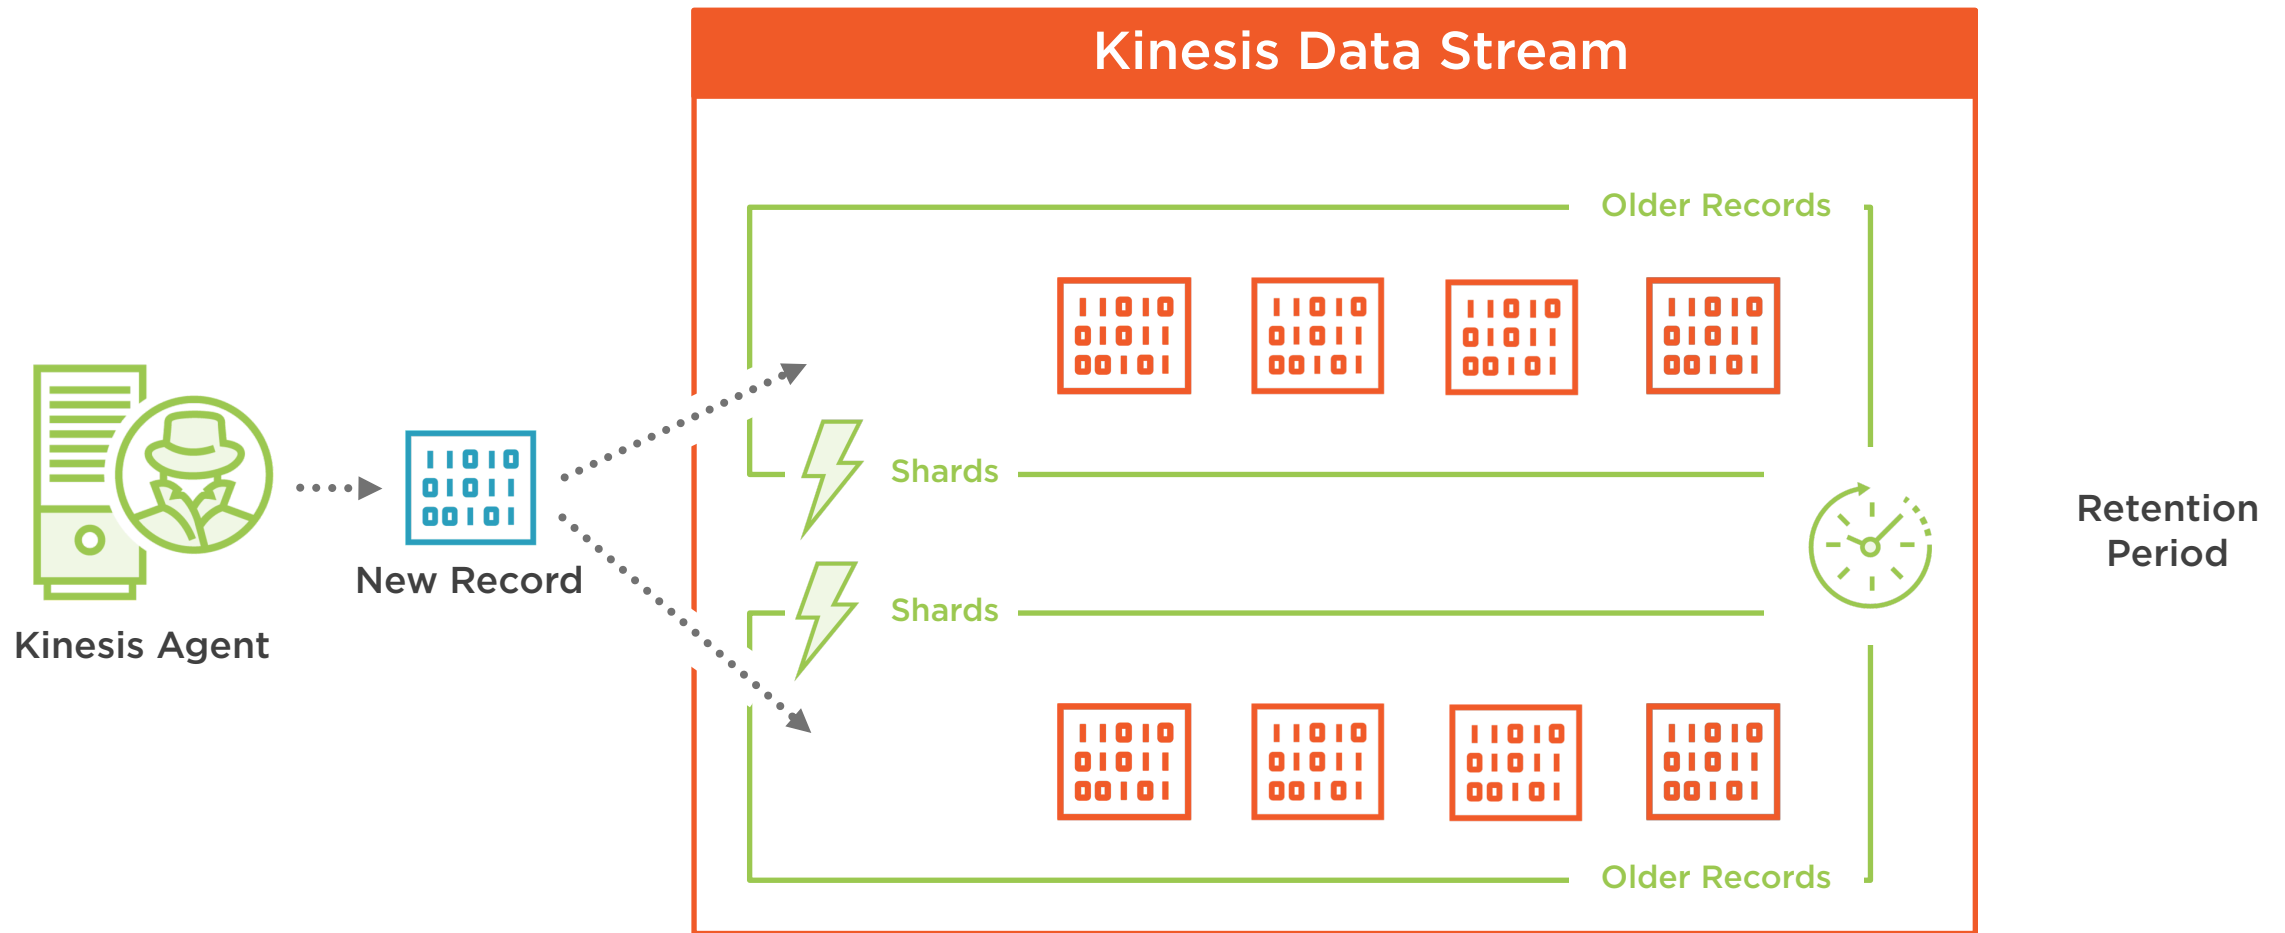
What's in a Kinesis Data Stream



What's in a Kinesis Data Stream



What's in a Kinesis Data Stream




What's in a Data Record?



Sequence Number




Partition Key



Data Blob



Stream and Record Limitations



Record data blob size limit - 1 MB



Number of streams per account – No limit



Soft limit of shards per account - 200 or 500 depending on region



Data retention length - 24 hours to 7 days



Shard Limitations

Ingest 1 MB
per second

Ingest 1000 records
per second

Read 2 MB per second

5 read transactions
per second

10 MB or 10k records
per read transaction



Resharding : Merging or Splitting a Shard



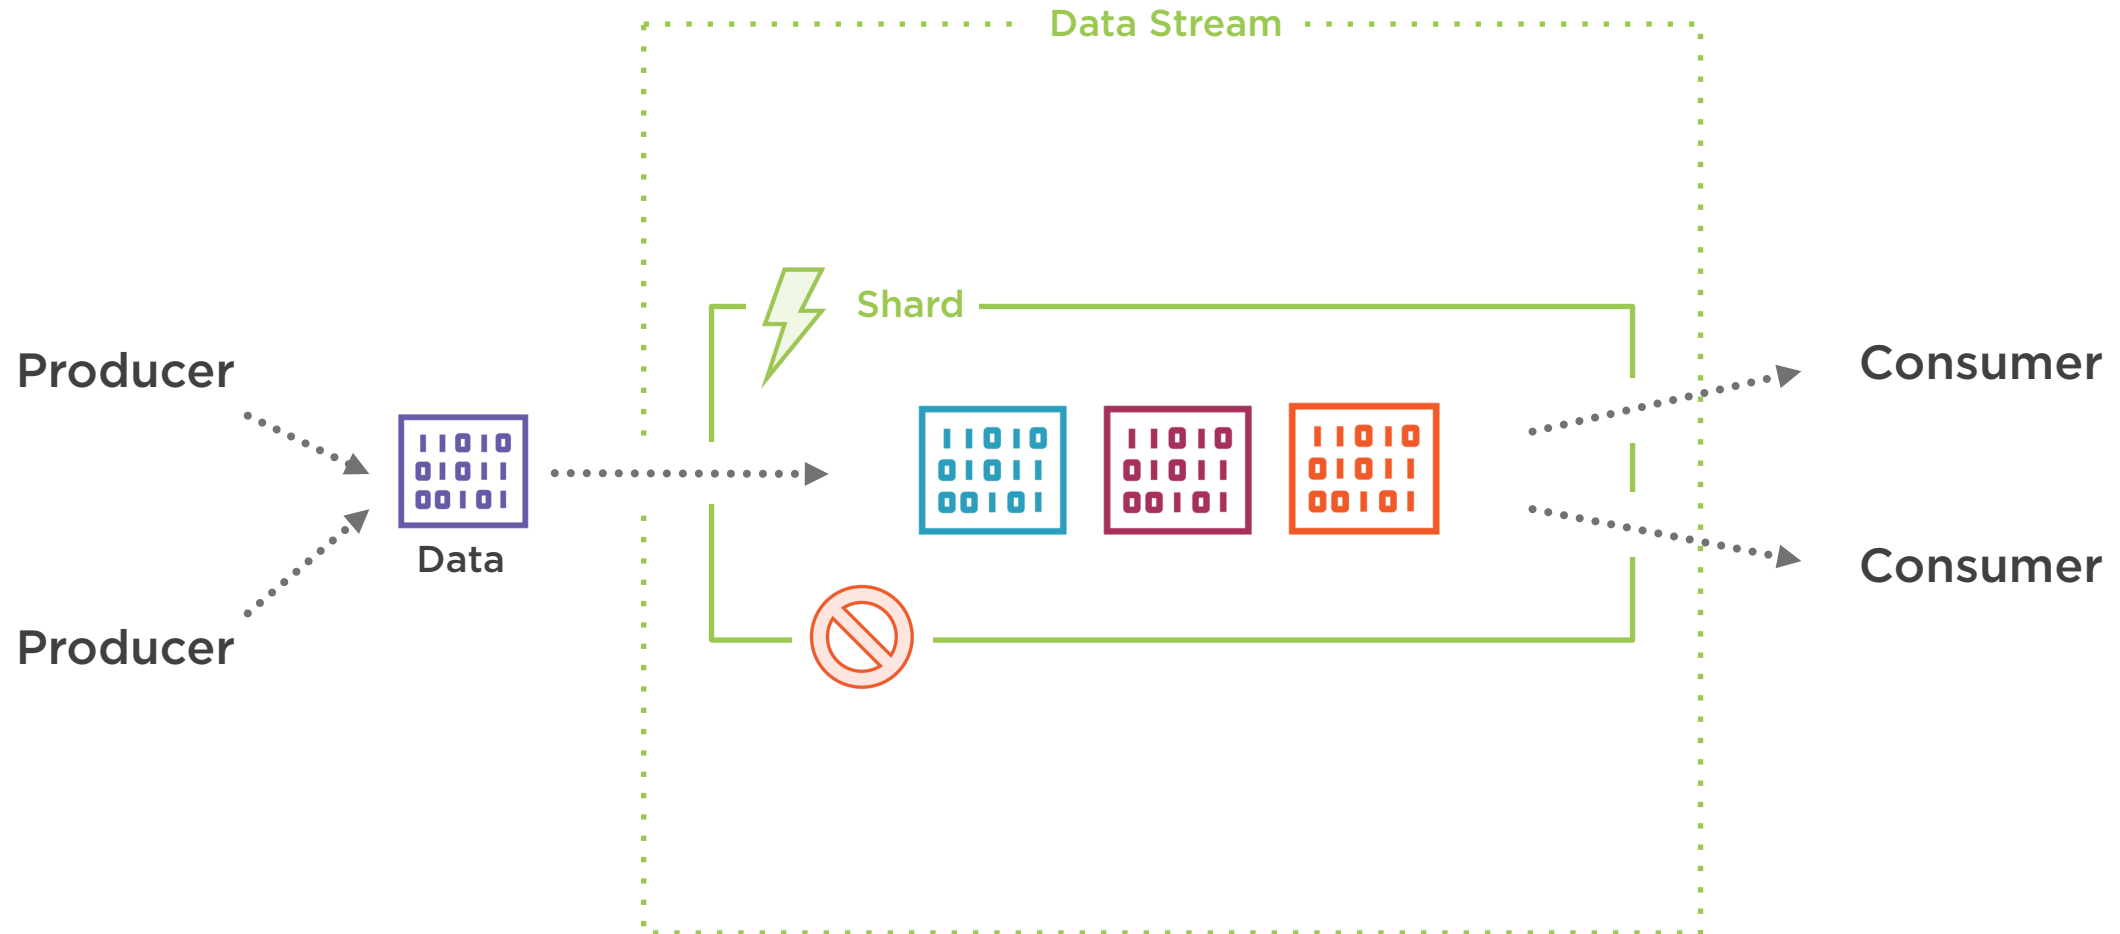
Shard Split



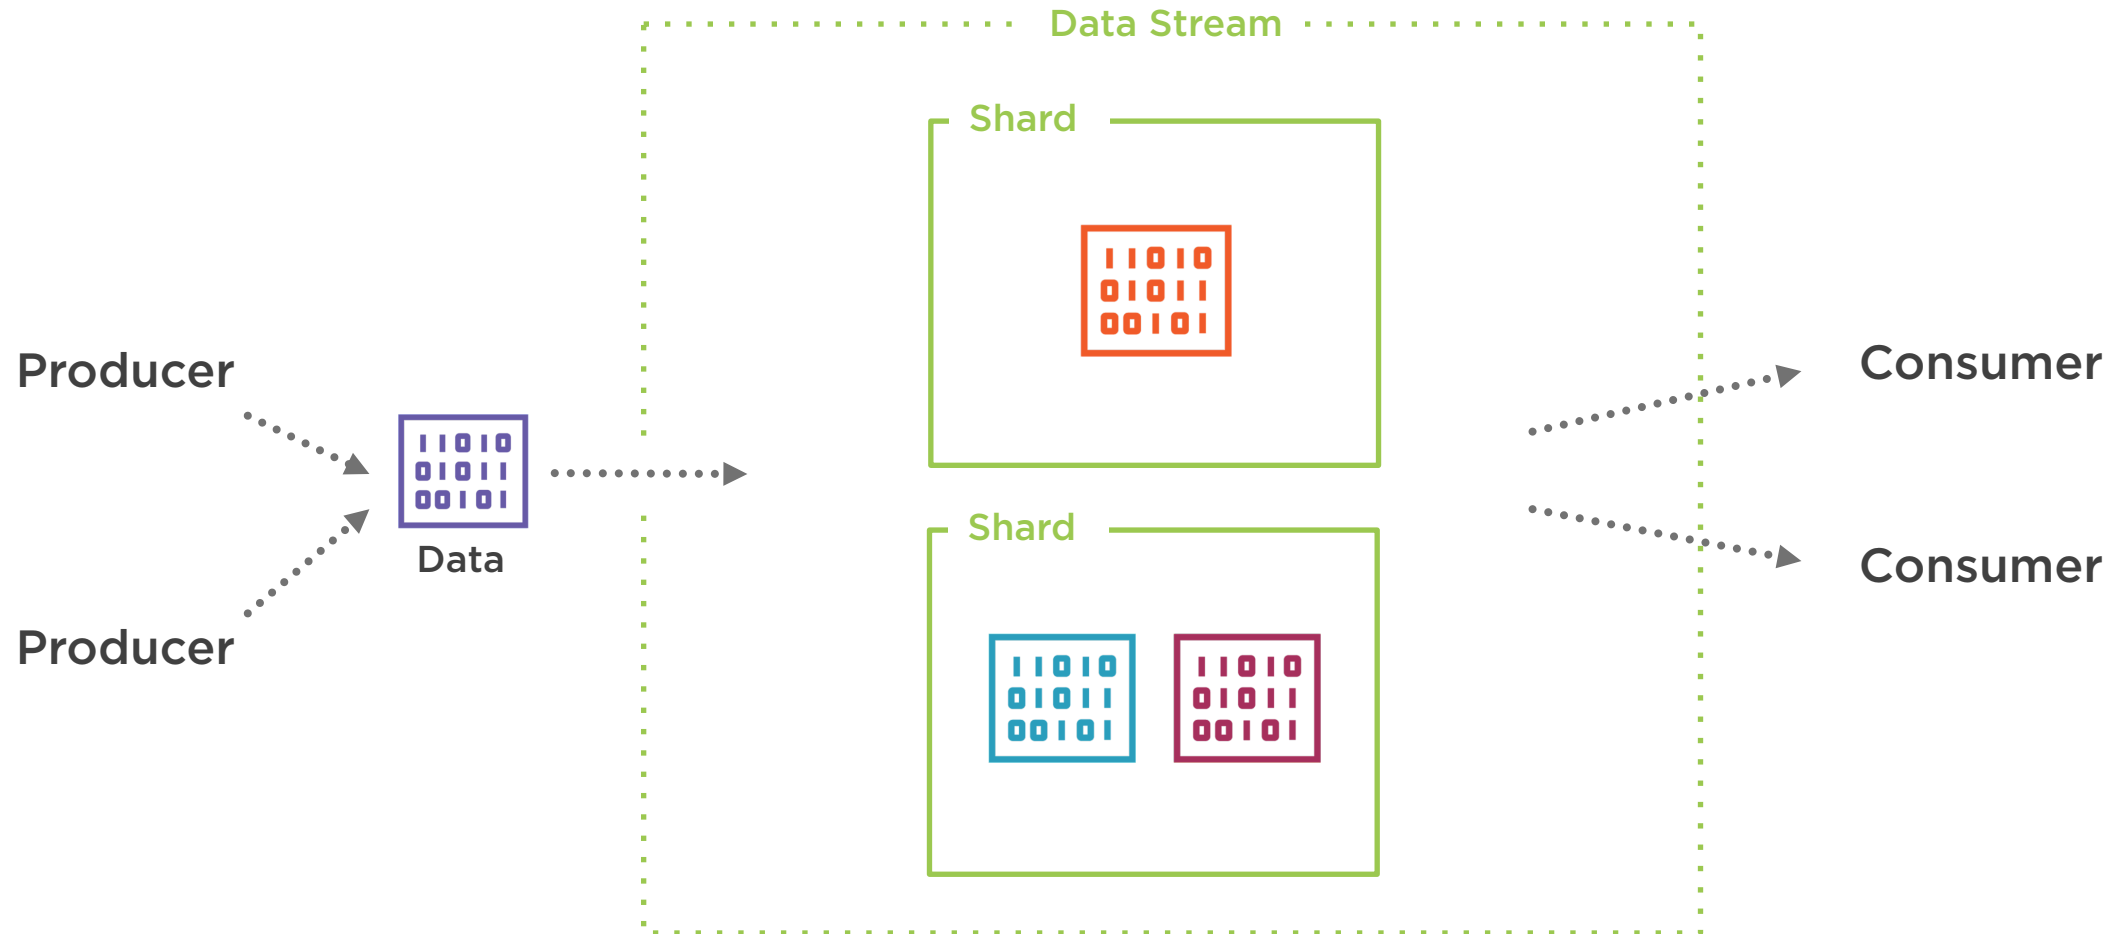
Shard Merge



Visualizing Kinesis Data Streams



Visualizing Kinesis Data Streams



Kinesis Data Streams Producers




Producers – Kinesis Agent



Kinesis Agent



Producer Library



AWS SDKs



What Is the Kinesis Agent?



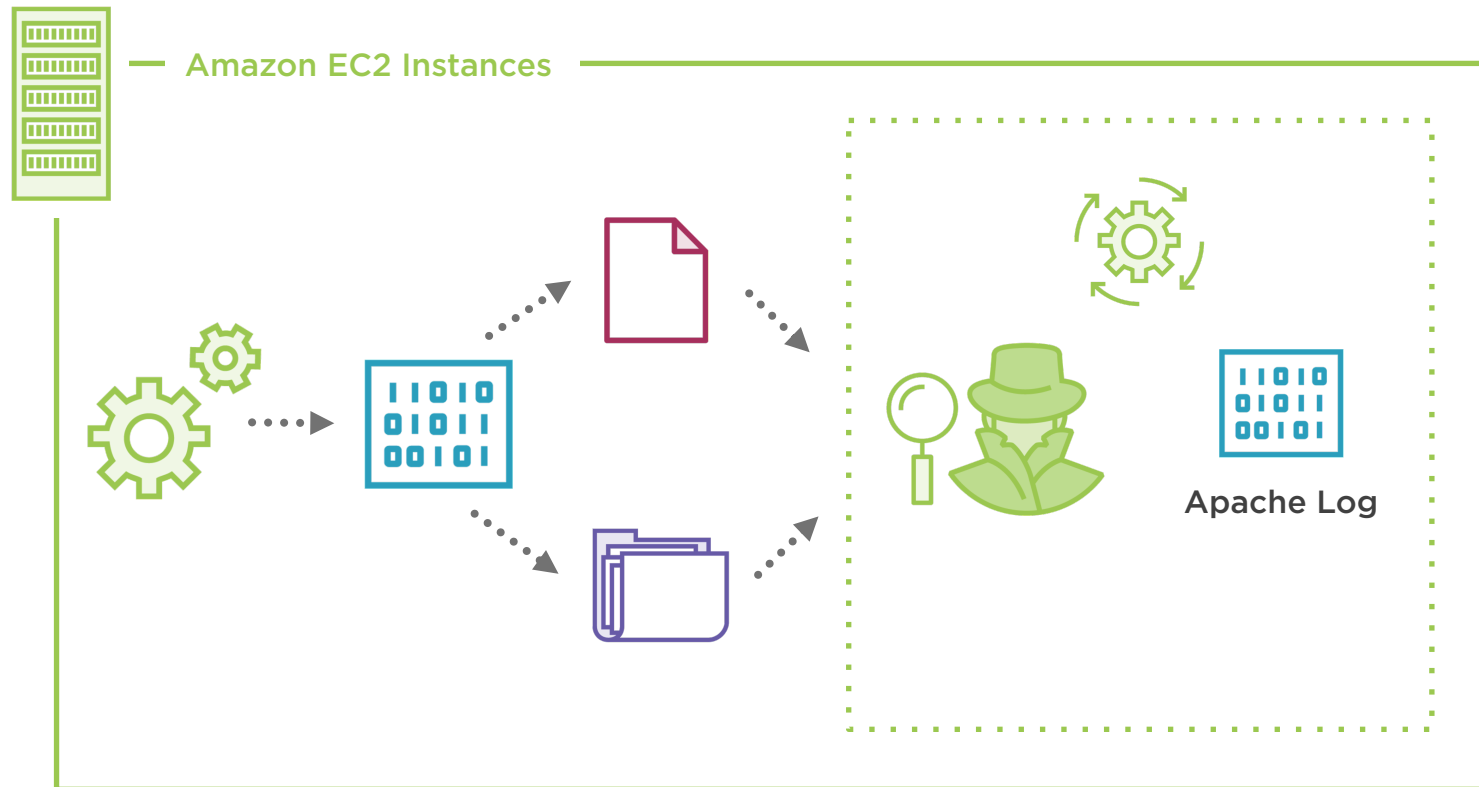
Java application

Streams logs from one
or many machines

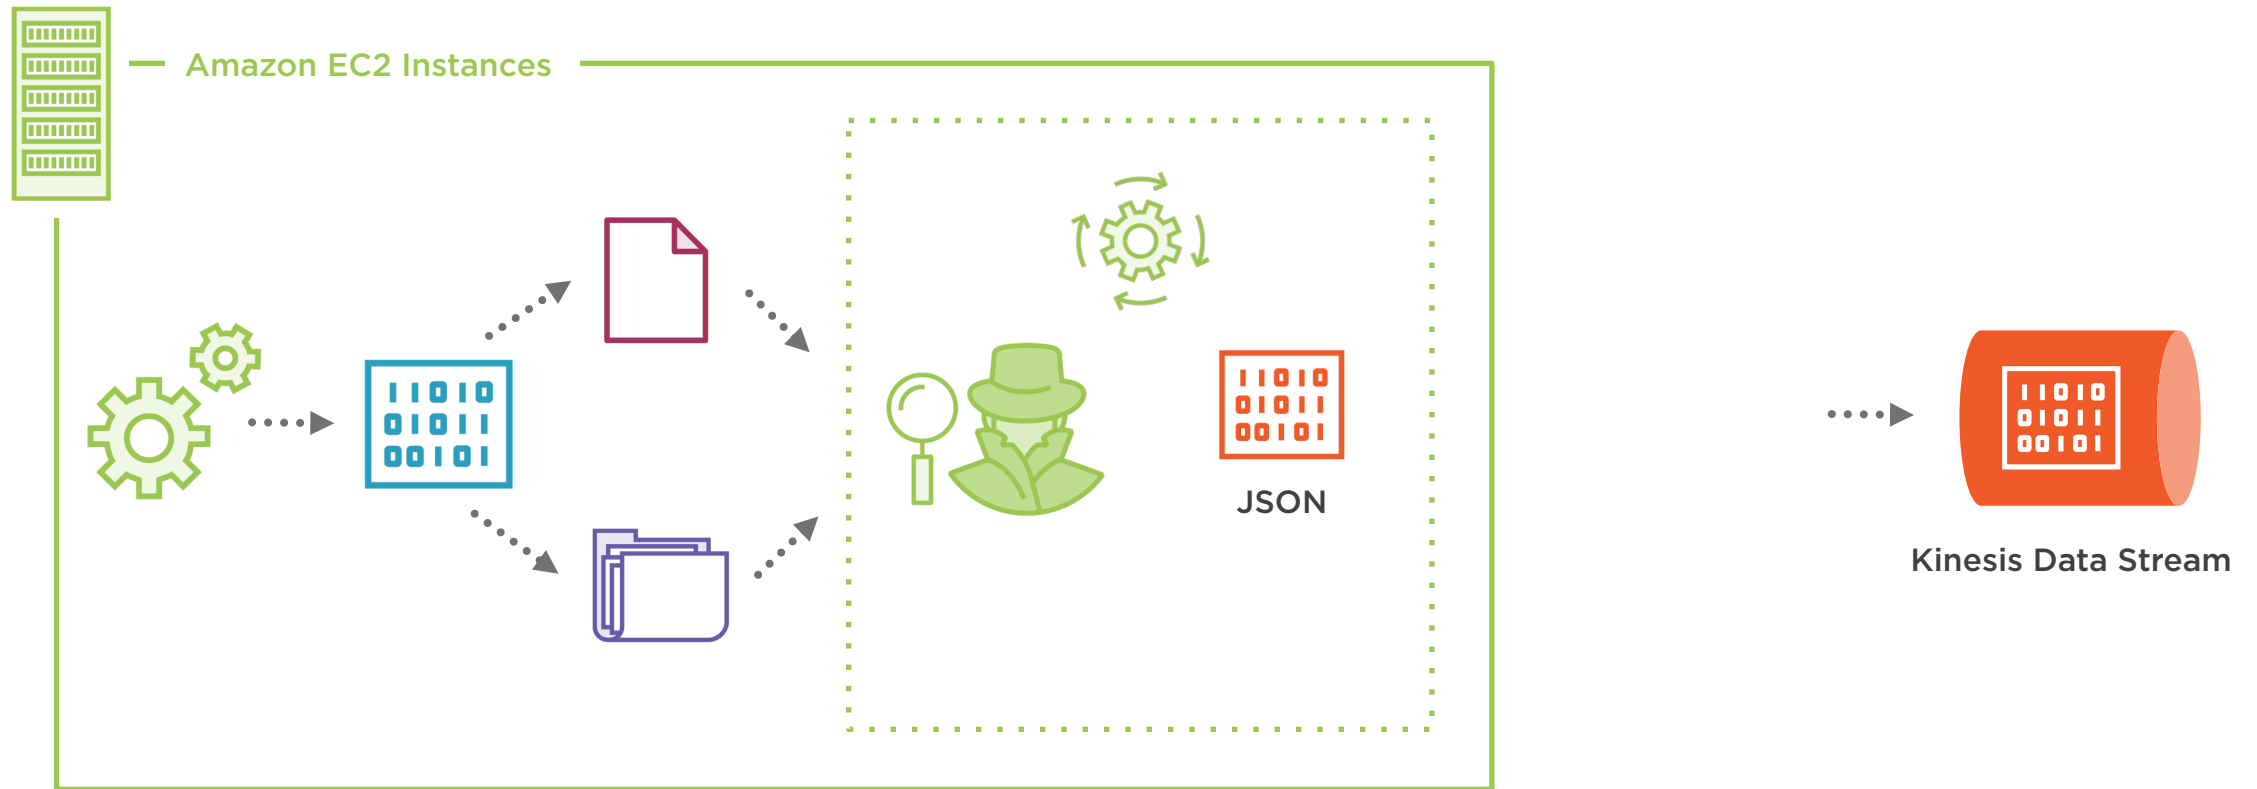
Lots of useful features
and integrations



Visualizing the Kinesis Agent in Action



Visualizing the Kinesis Agent in Action



Kinesis Agent Log Formats

**Multi-line to
Single Line**

**CSV/PSV/DSV
to JSON**

**Apache Log or Syslog
formats to JSON**



Why Use the Kinesis Agent?



Performance benefits

Lots of built-in functionality

CloudWatch integration


Producers – Producer Library



Kinesis Agent



Producer Library



AWS SDKs



What Is the Producer Library?



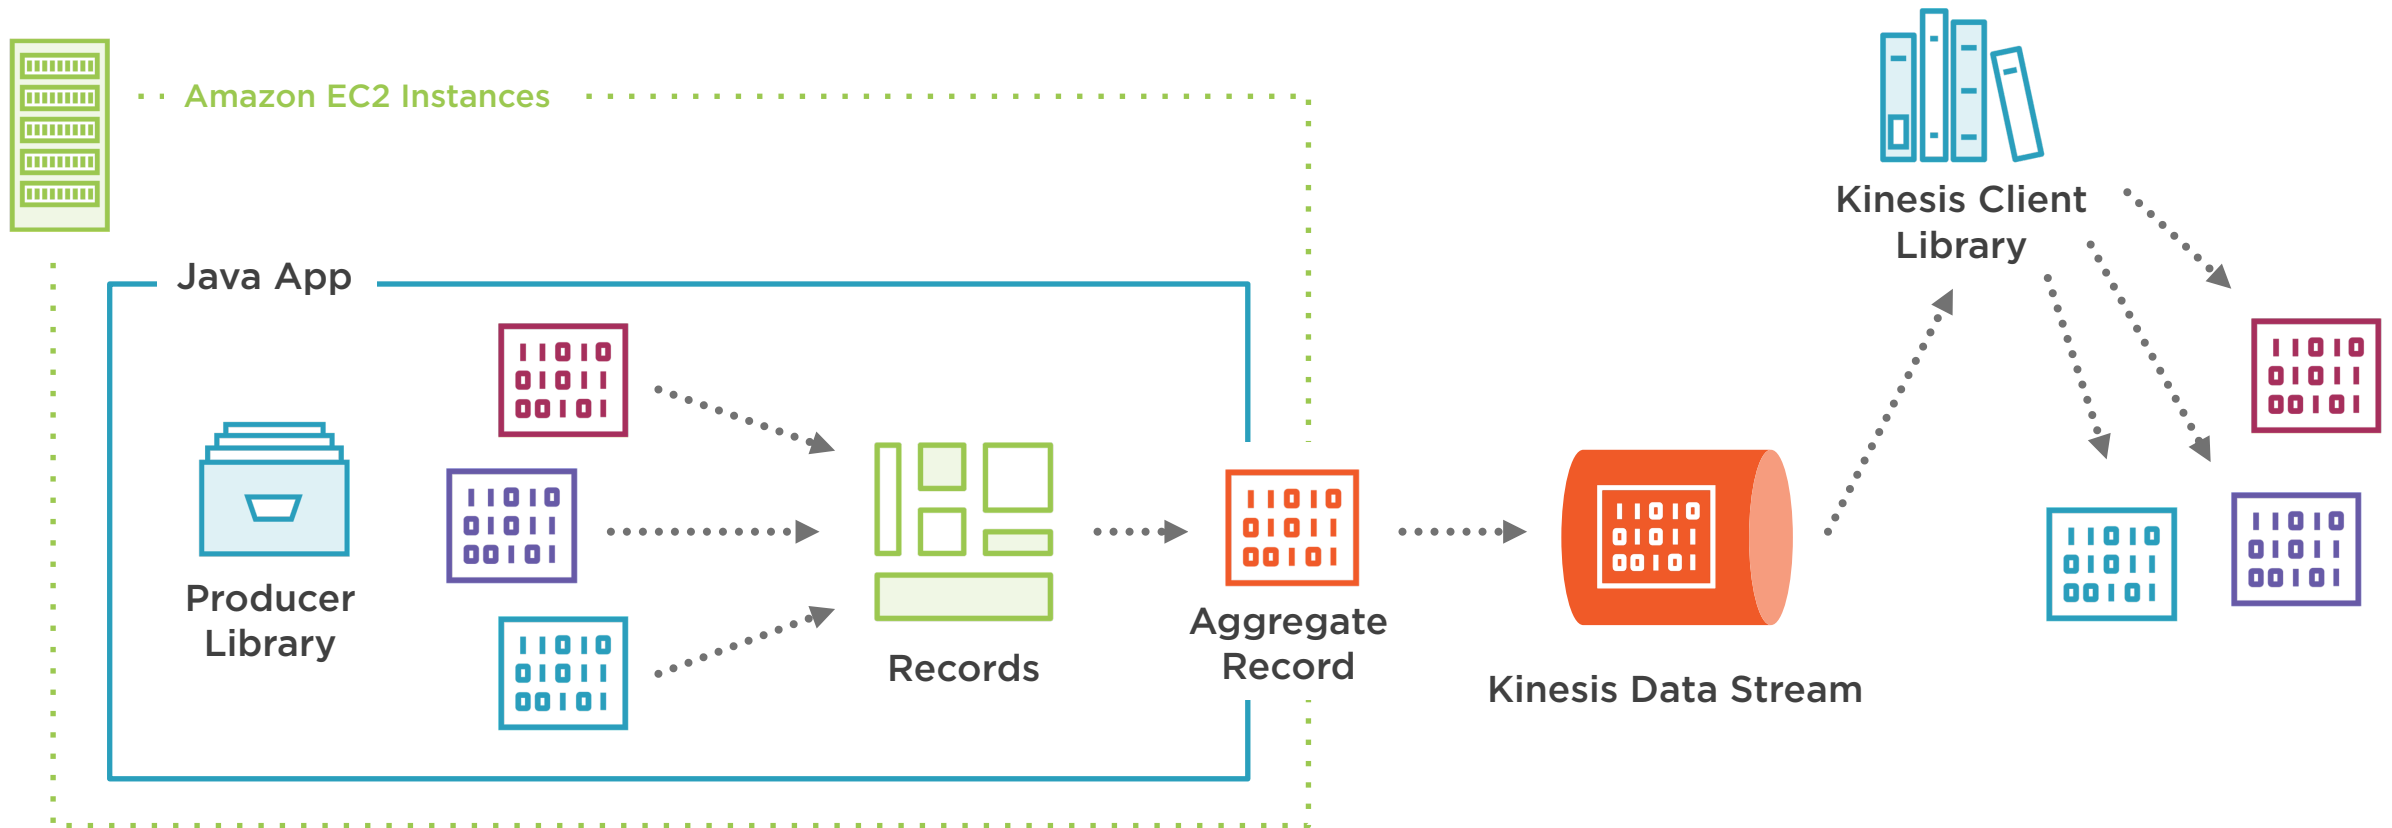
Library for Java apps

**Helps collect and
aggregate logs
to Kinesis**

**Useful when working
with high-rate producers**



Visualizing the Producer Library



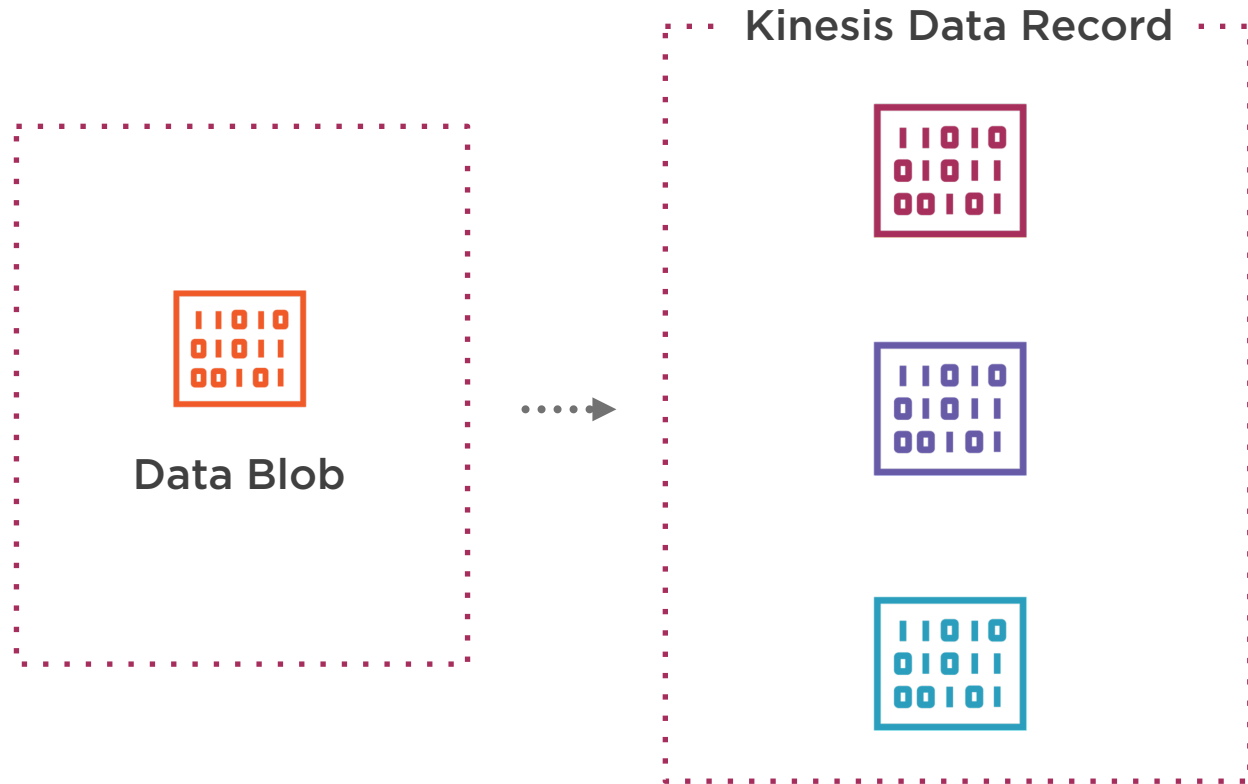
Kinesis Producer Library - Batching

Aggregation

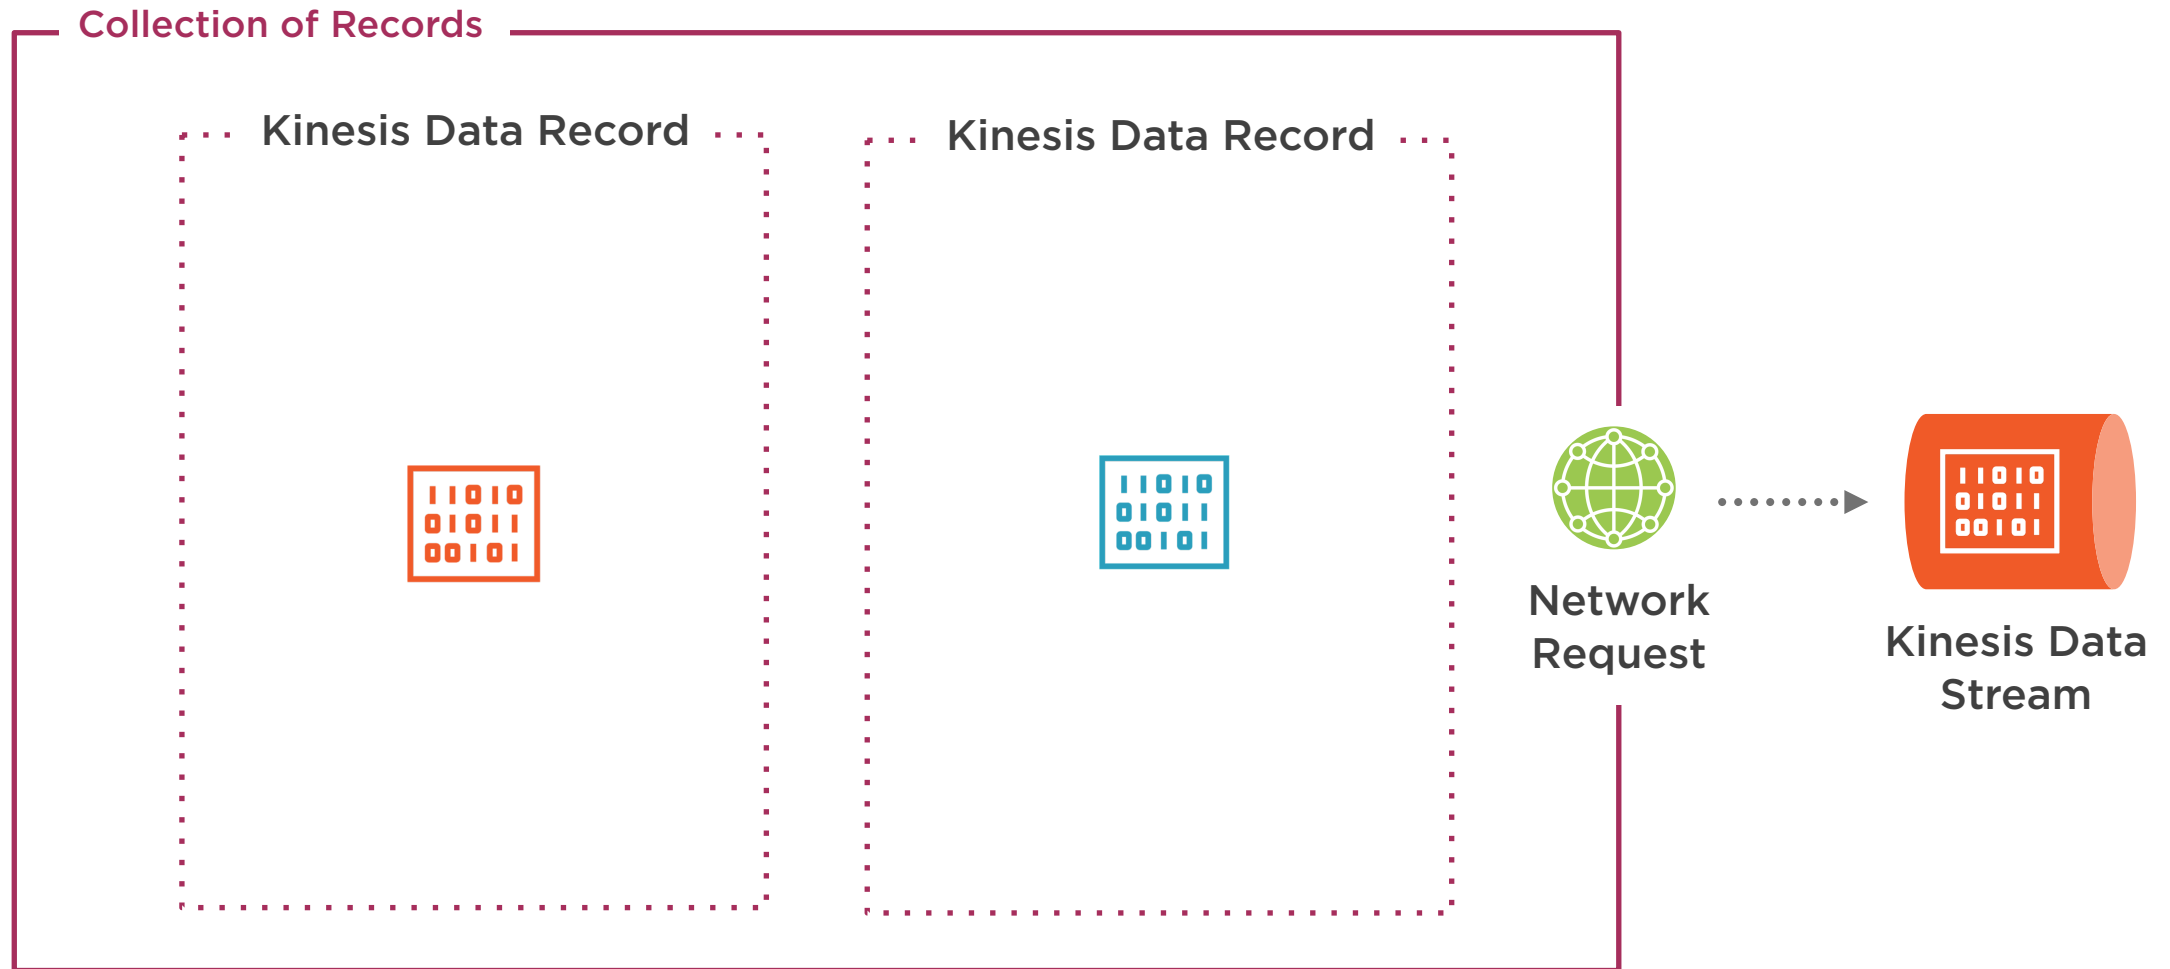
Collection



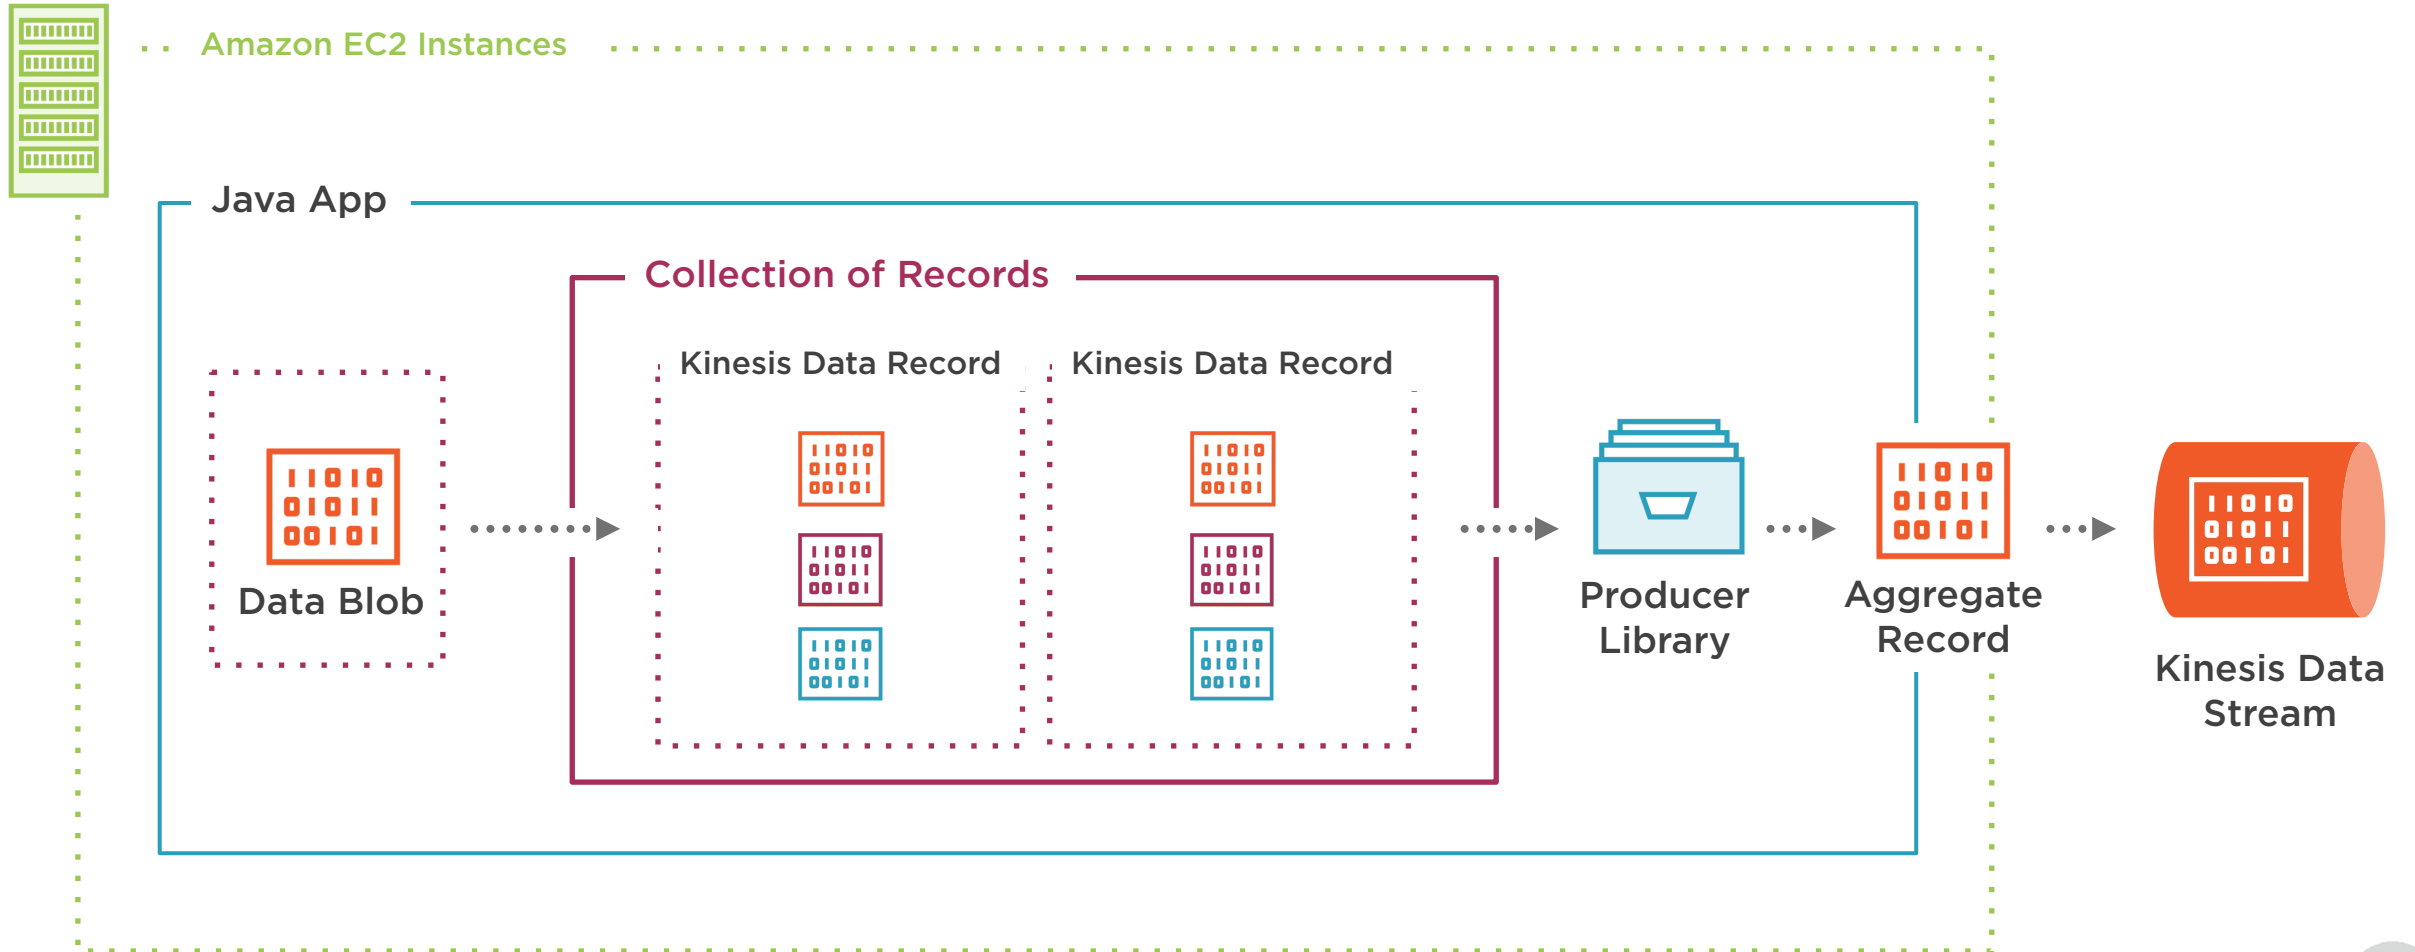
Batching - Aggregation



Batching - Collection



Batching - Collection and Aggregation



Why Use the Producer Library?

**Helps handle errors with
retry logic**

**Batching and performance with
high-rate producers**

CloudWatch integration

**Integrates with the Kinesis
Client Libraries**




Producers – AWS SDKs



Kinesis Agent




Producer Library



AWS SDKs



Why Use the AWS SDKs?



Supports multiple
programming languages

Embed inside
mobile apps

Low rate producers



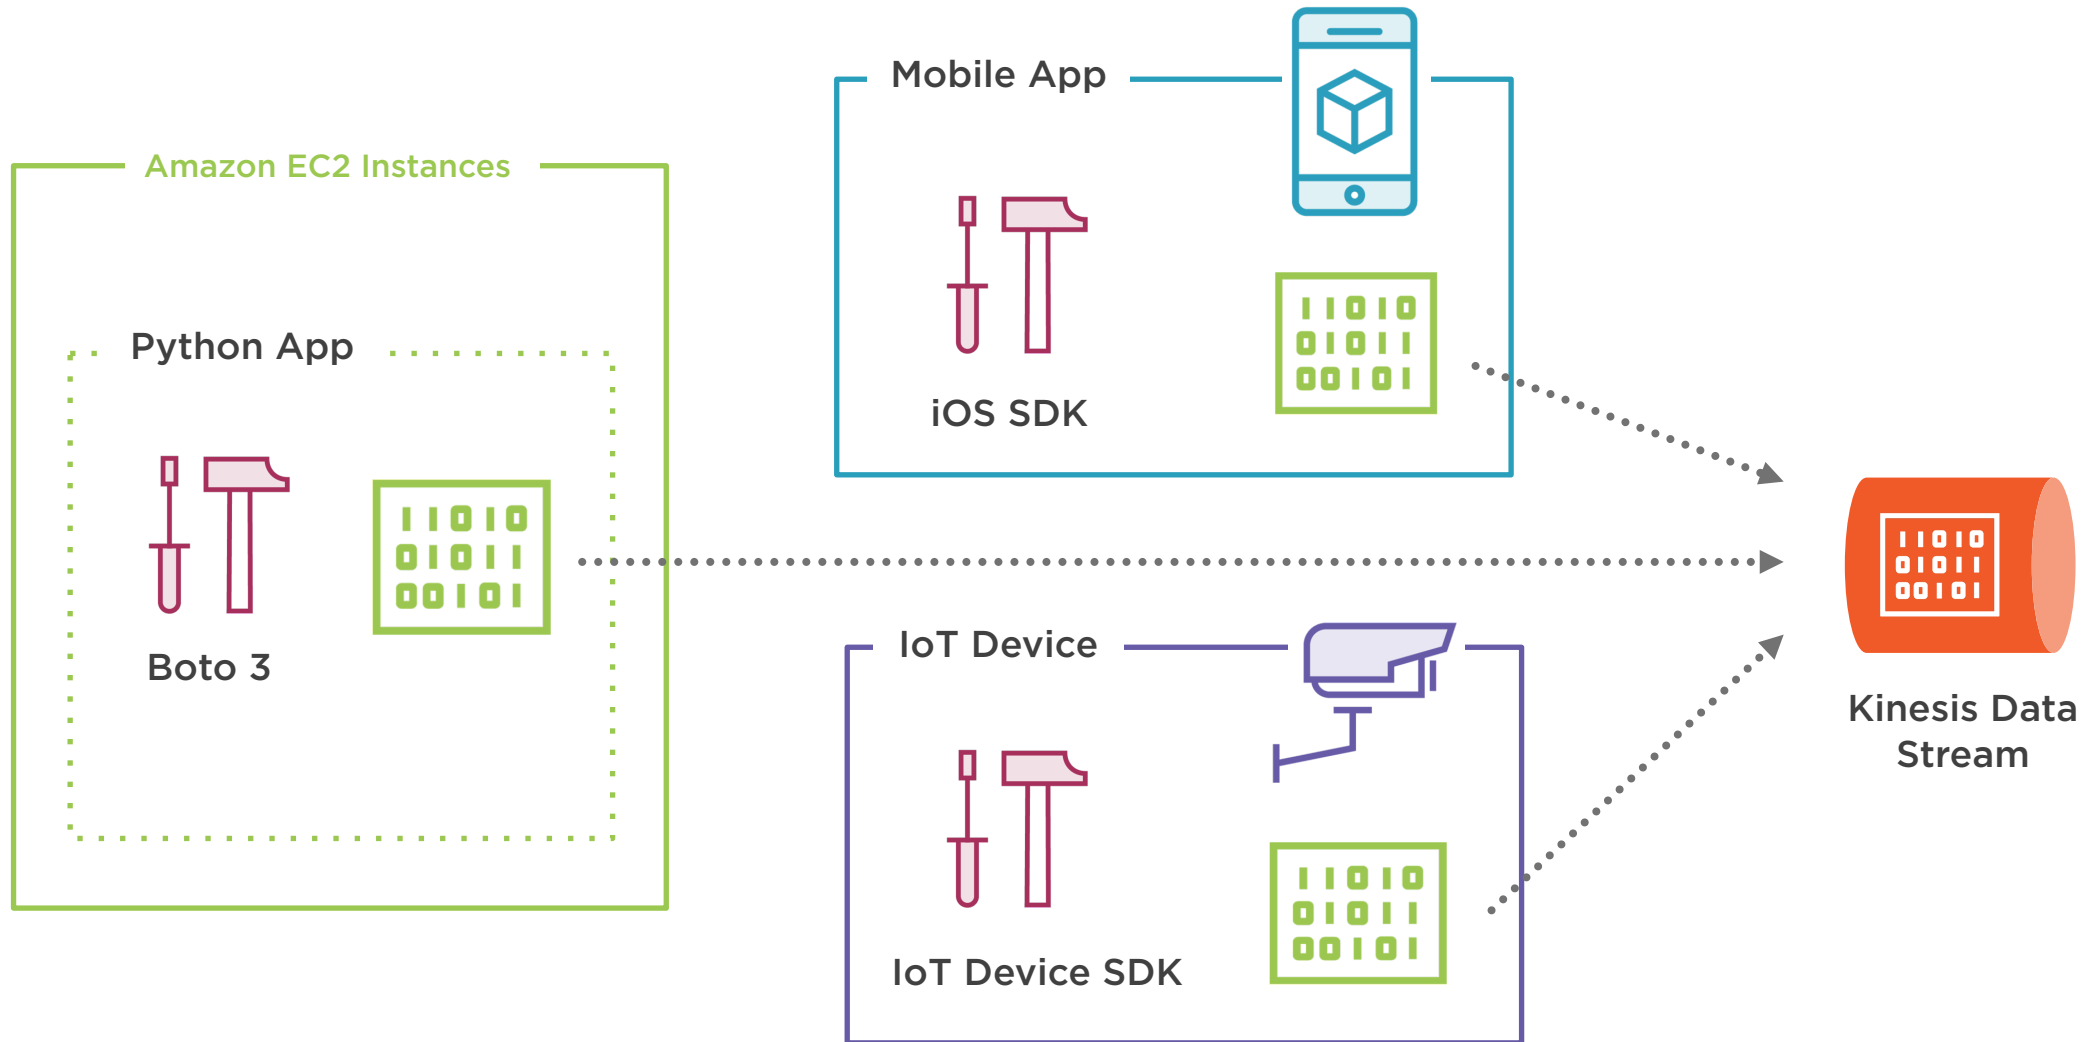
Options for the AWS SDKs

PutRecord

PutRecords



Visualizing the AWS SDKs




Kinesis Data Streams Consumers



Consumers – Client Libraries



Client Libraries



AWS SDKs



AWS Services



What Are the Kinesis Client Libraries?



Tool for processing data
from Kinesis Streams

Variety of programming
languages supported

KCLs help with common
use cases



Kinesis Client Libraries – Supported Languages

Ruby

.Net

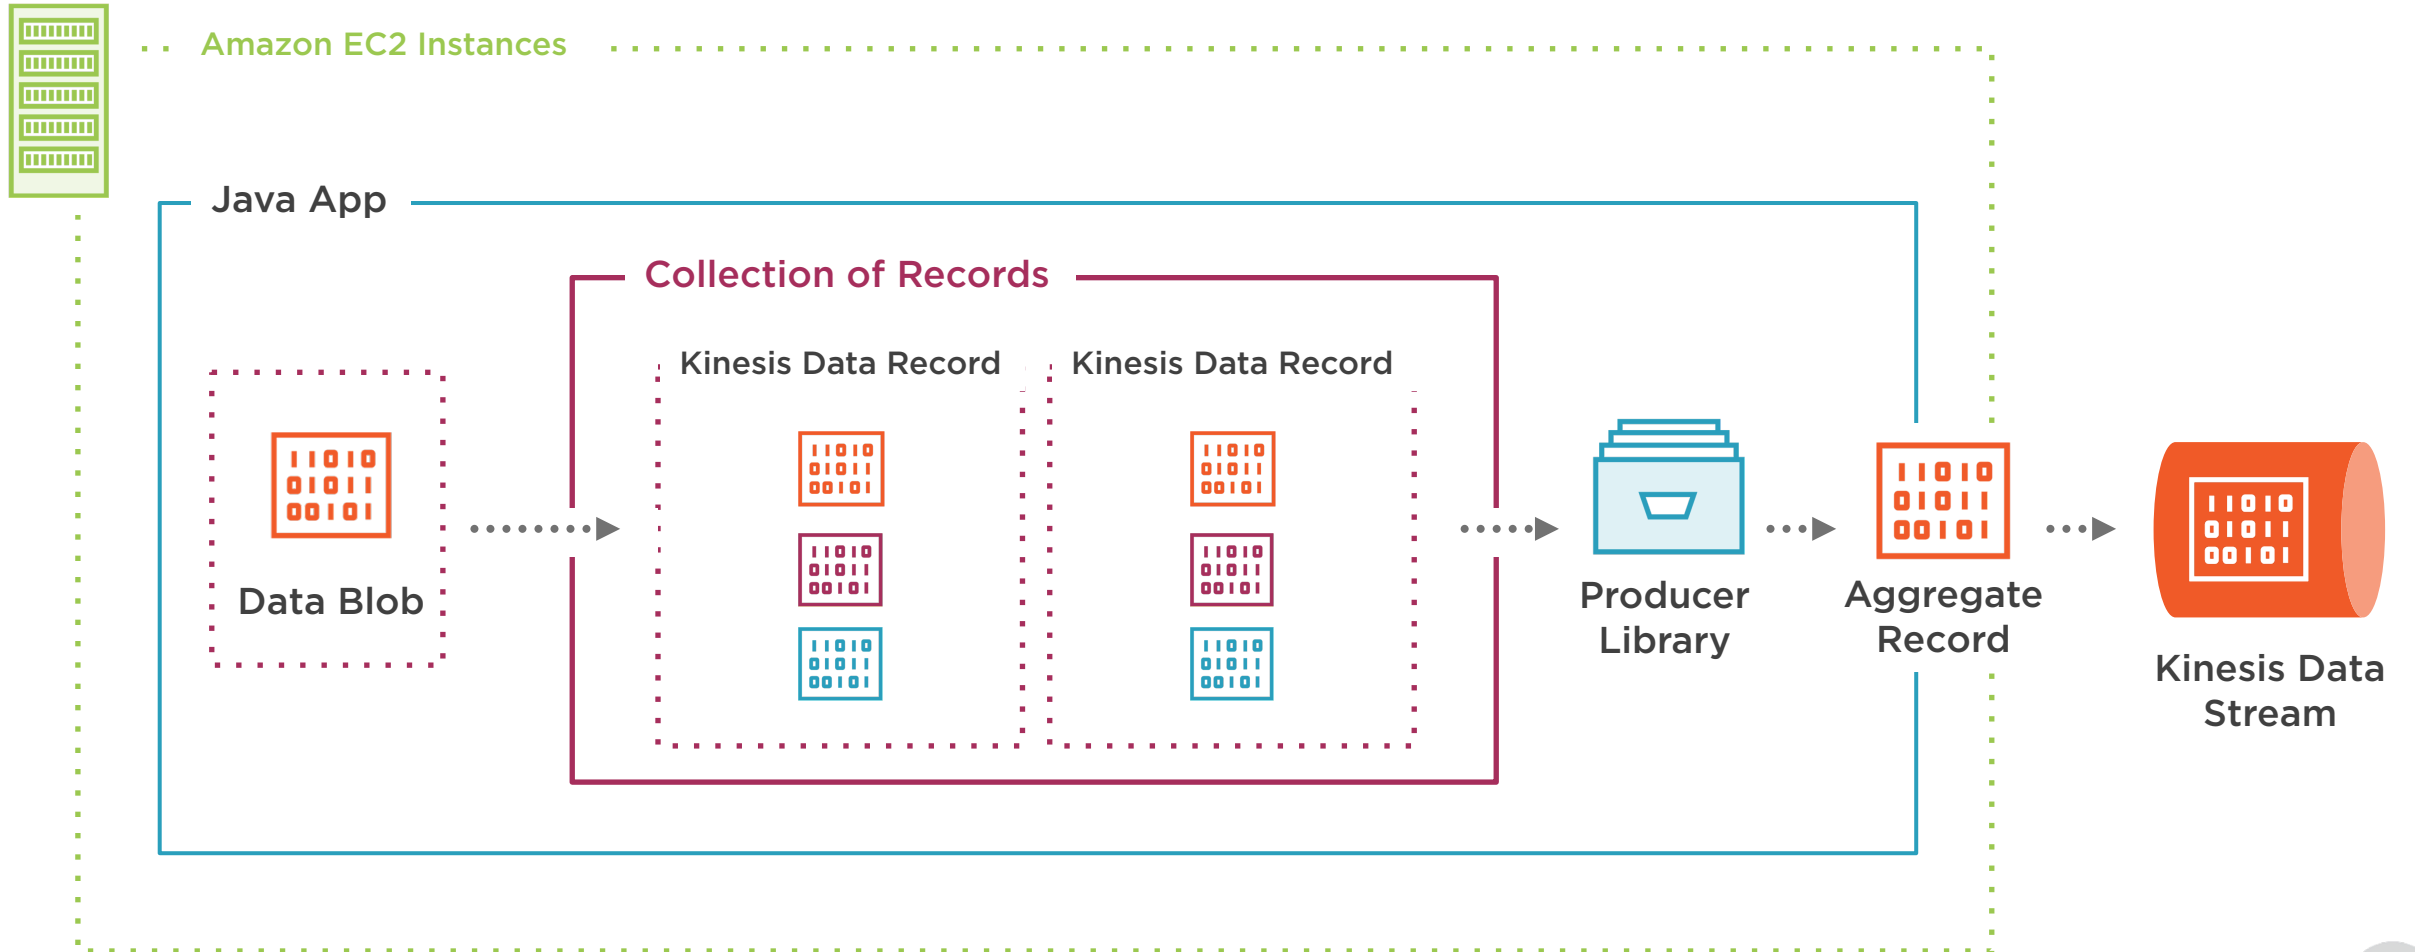
Node.js

Java

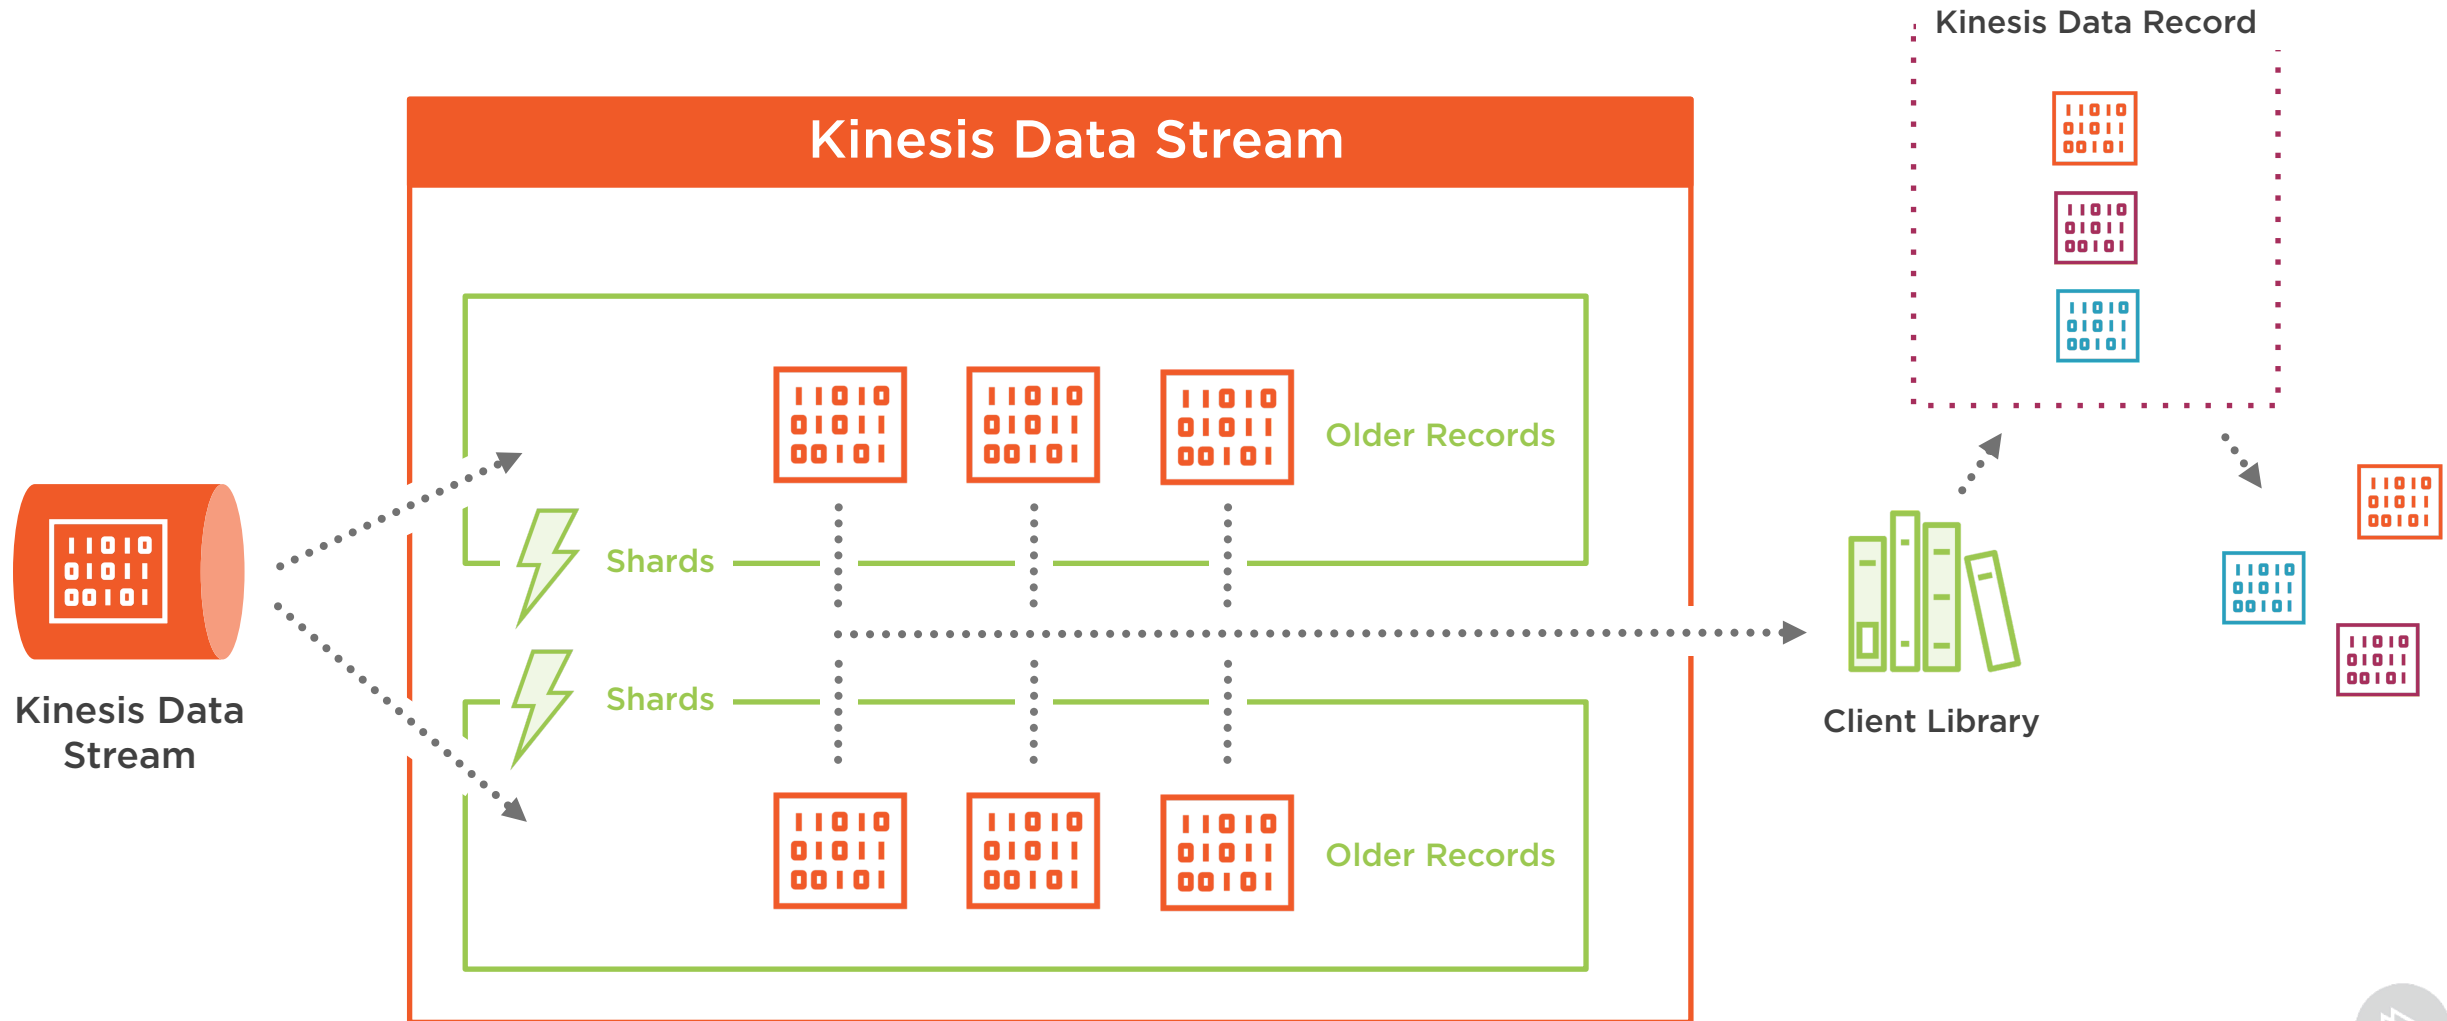
Python



Visualizing the Kinesis Client Library



Visualizing the Kinesis Client Library



Throughput Options



Shared throughput




**Dedicated throughput
(Enhanced Fan-out)**

Consumers – AWS SDKs



Client Libraries




AWS SDKs



AWS Services



Why Use the AWS SDKs?



Control over specific
API actions

Can integrate with any
SDK programming
language


Run apps in EC2, AWS
Lambda, containers
or elsewhere



Consumers – AWS Services



Client Libraries



AWS SDKs



AWS Services



Consumers – AWS Services



AWS Lambda



Kinesis Data Analytics



Kinesis Delivery
Firehose



AWS Lambda Consumer



“Serverless” compute




Runs in response to new records



Records are batched and send to Lambda



AWS Lambda Consumer



Kinesis Data Analytics

**Stream data to Kinesis
Data Analytics**

Run SQL queries

**Send output to
other streams**



Kinesis Data Firehose



Data flows from Kinesis Data Stream to Firehose



Firehose can process and transform data



Delivery to other destinations (S3, ElasticSearch, others)



Demo



WARNING

- Not in AWS free tier

Storyline

- Sales data from web API backend
- Searching for popular products
- Live analytics of streaming data

Create Kinesis Data Stream

- AWS Console
- Send data in with Boto3

Create Kinesis Data Analytics Application

- AWS Console
- Writing SQL queries



Summary



Kinesis Data Streams concepts

- Producers
- Shards and records
- Consumers
- Shared and dedicated throughput

Hands-on with:

- Kinesis Data Streams
- Kinesis Data Analytics





Up next:
Amazon Kinesis Data Firehose

