

Amazon Kinesis Data Firehose



Fernando Medina Corey

CLOUD ARCHITECT

@fmc_sea www.fernandomc.com



Overview



Basic Concepts

- Records, buffers, destinations, retention
- Firehose Limits

Data Transformations with Firehose

- Visualizing transformations
- Lambda transformation examples

Firehose Producers

Firehose Destinations

Demo: Creating a delivery stream


When to use Firehose vs. Data Streams



Kinesis Data Firehose: Basic Concepts



Records, Data Retention, and Destinations



Data Records




Data Retention
(24 hours max)



Destinations
(S3, Splunk, etc.)

Firehose Buffer Size and Interval

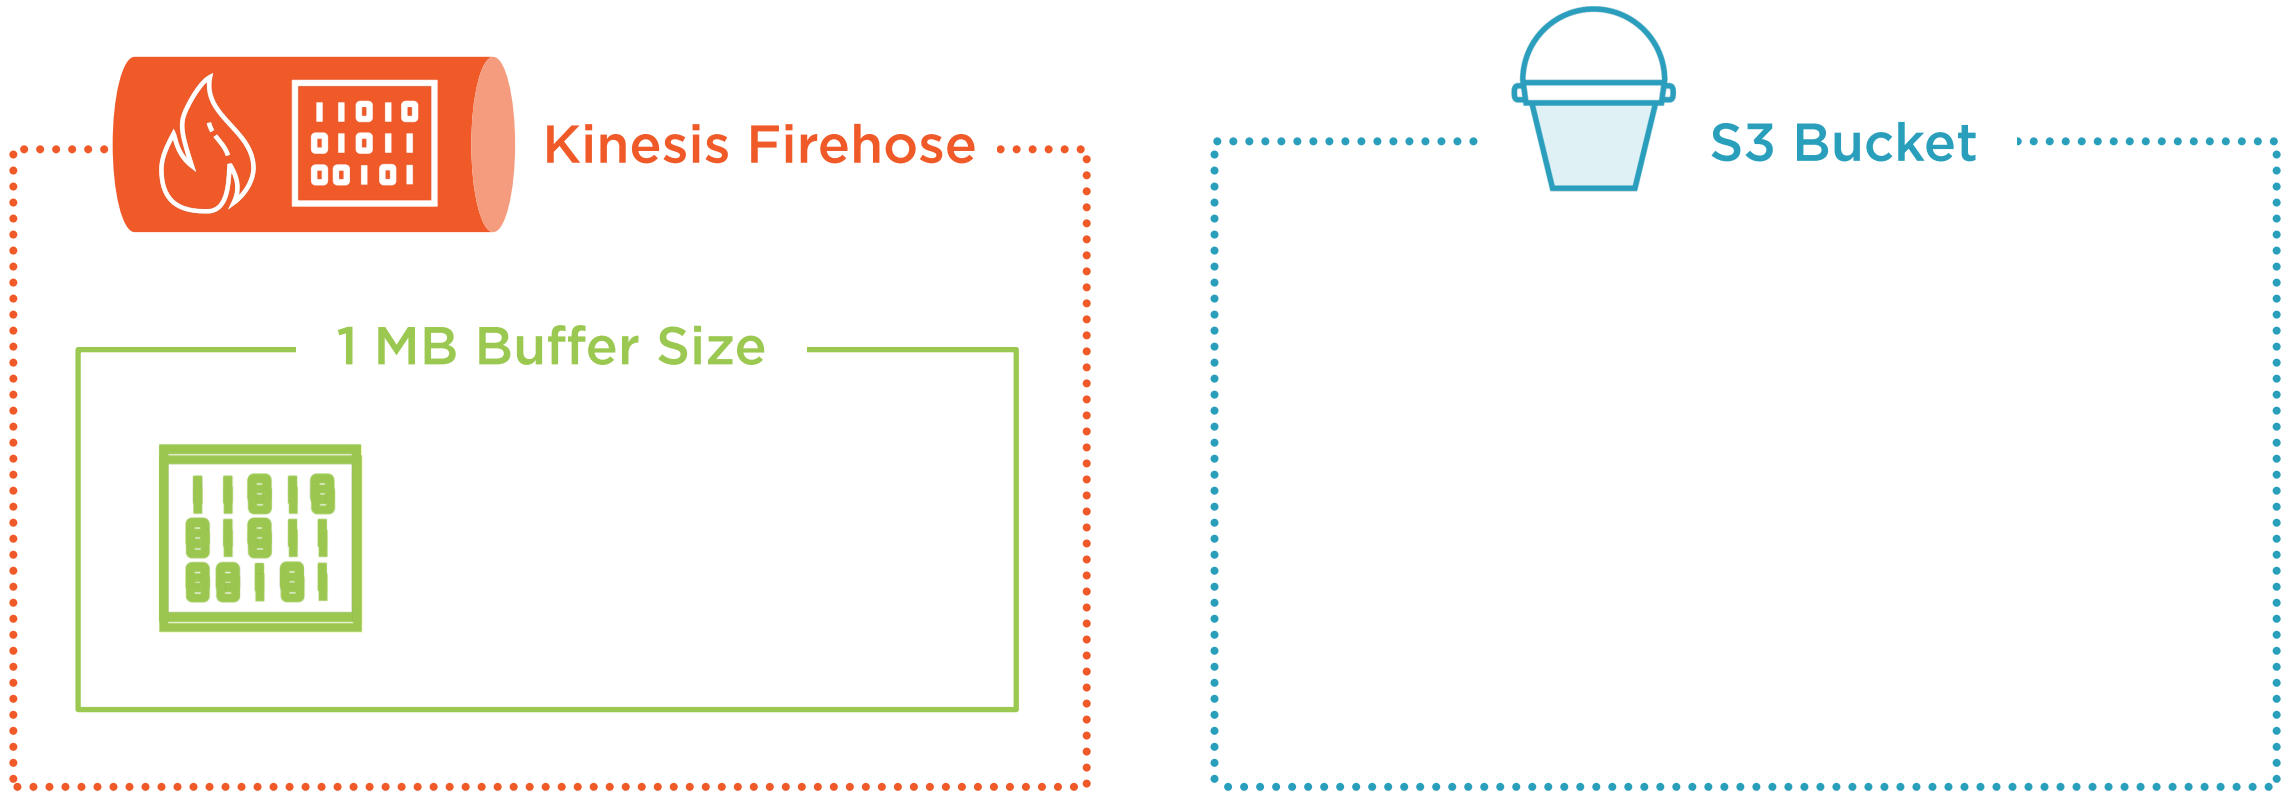


Buffer Size
(1MB – 128MB)

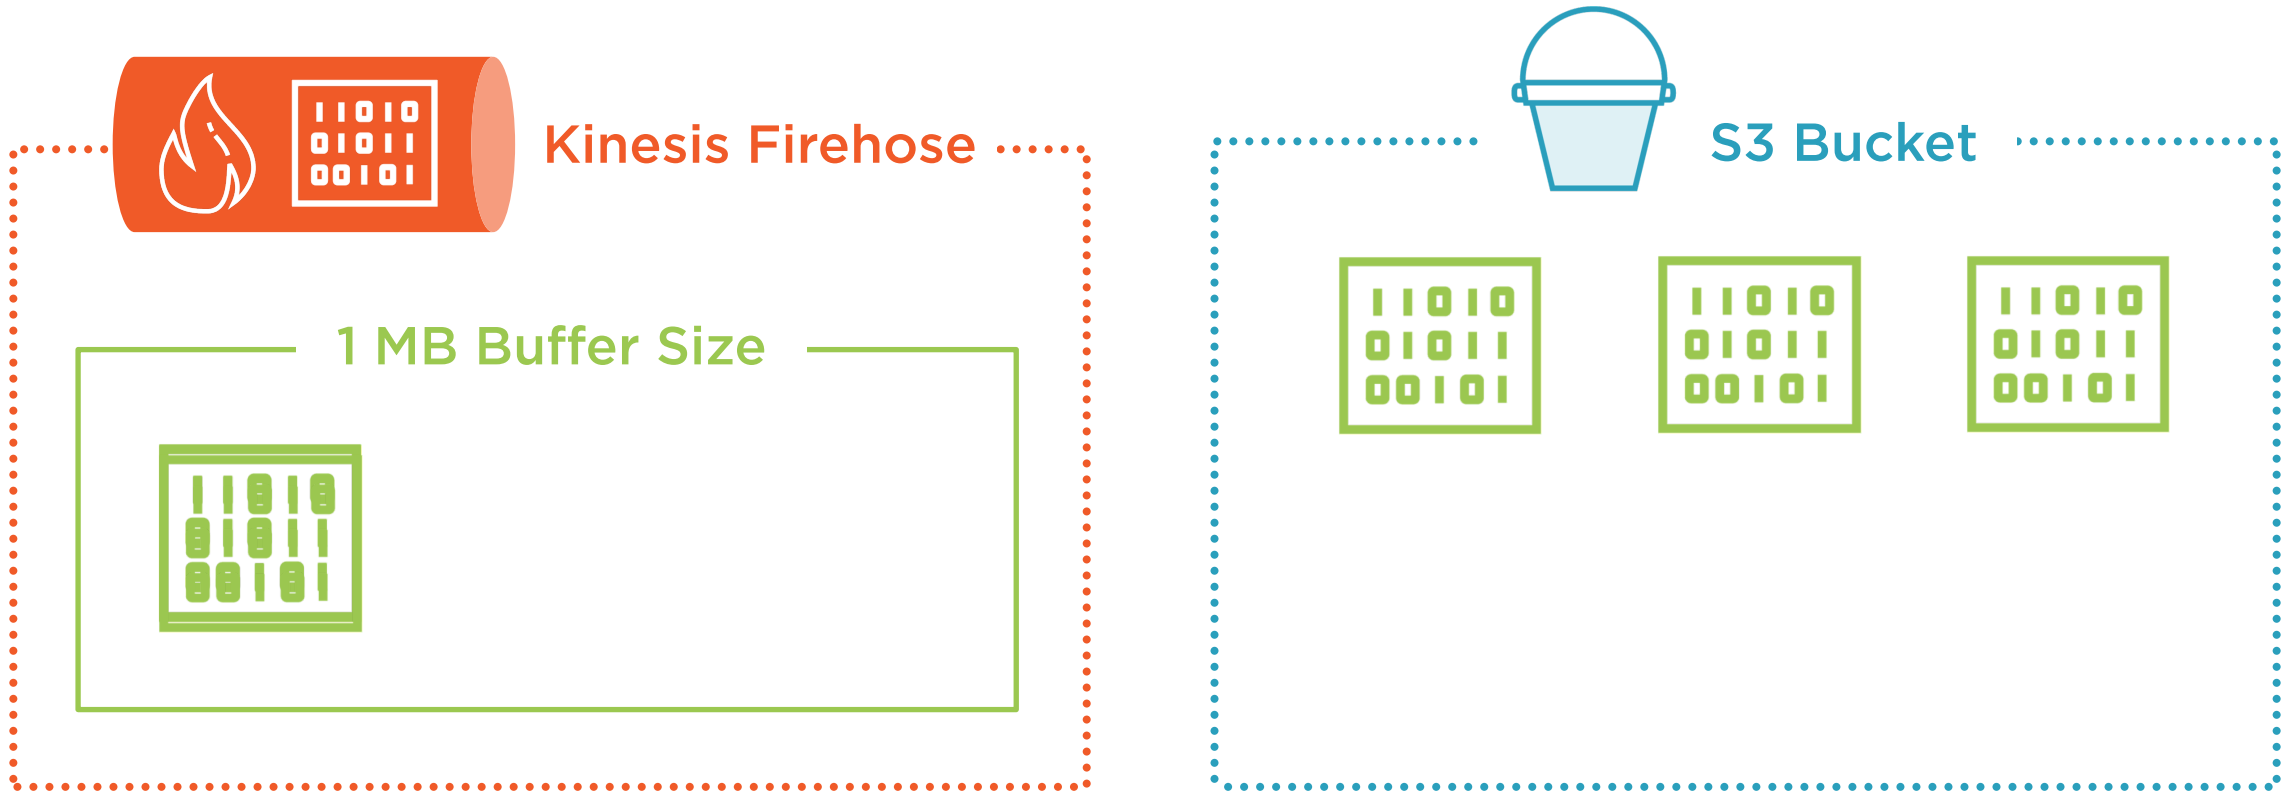


Buffer Interval
(60 – 900 seconds)

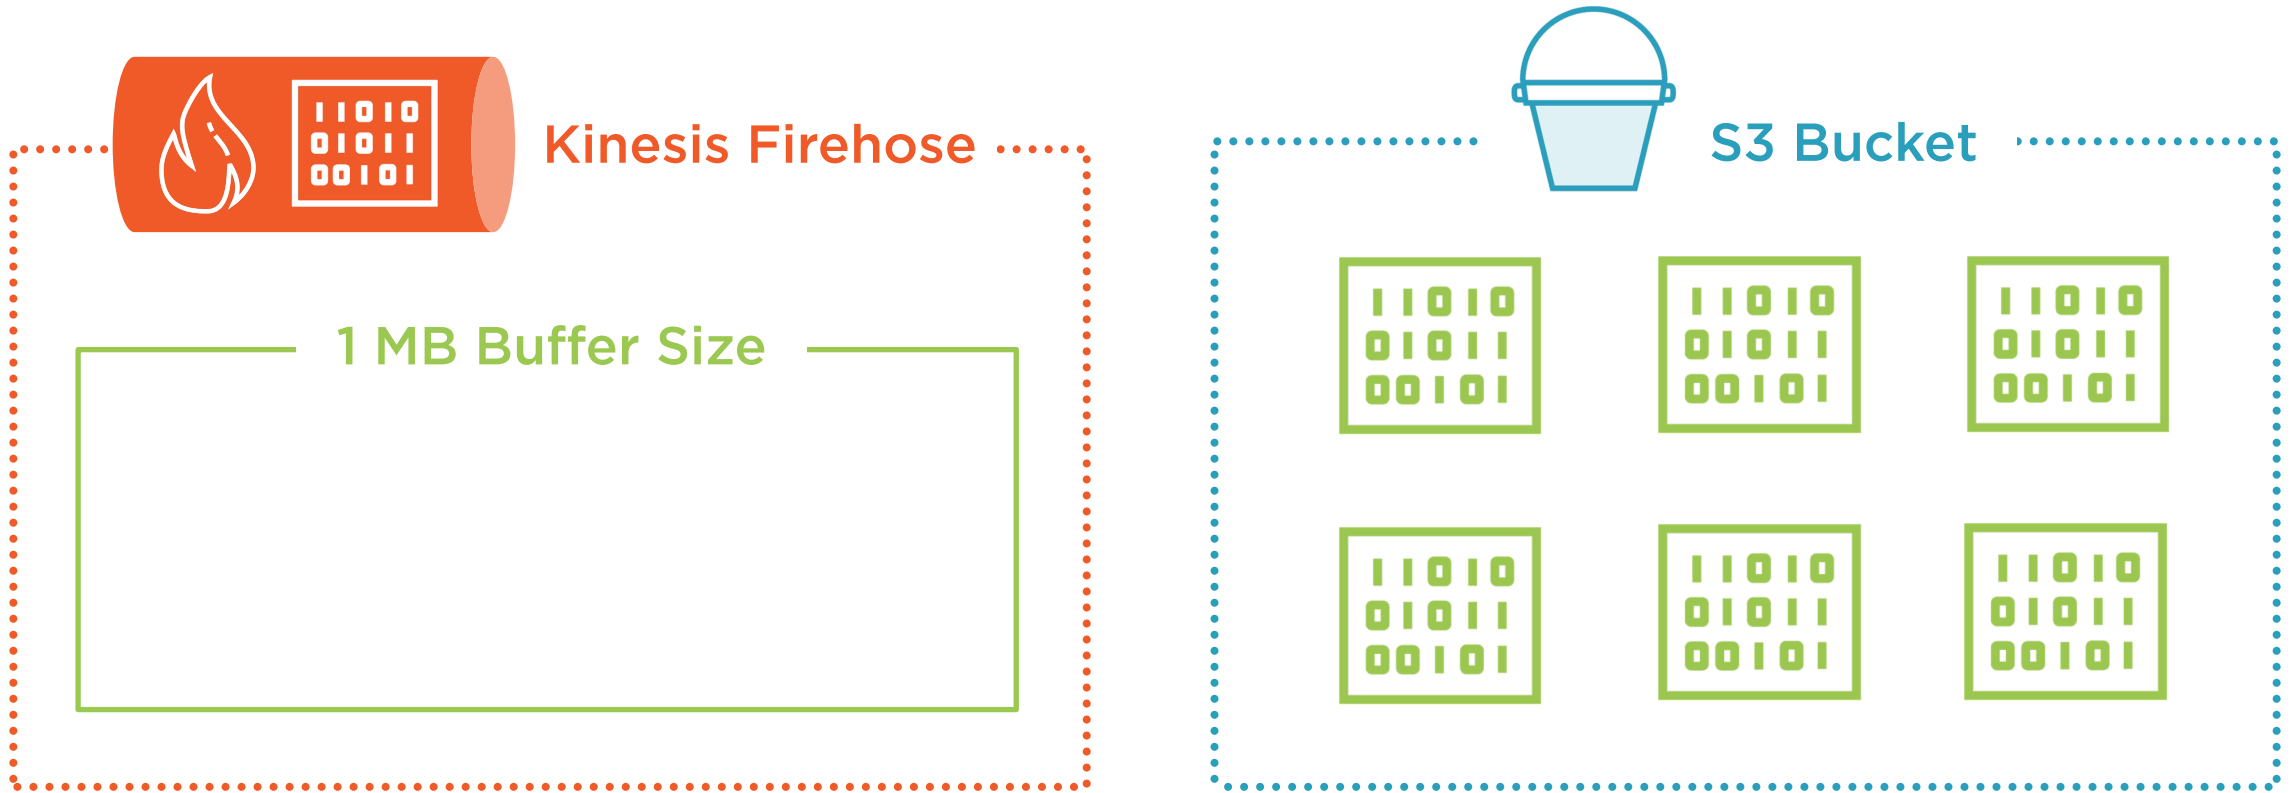
Buffer Size



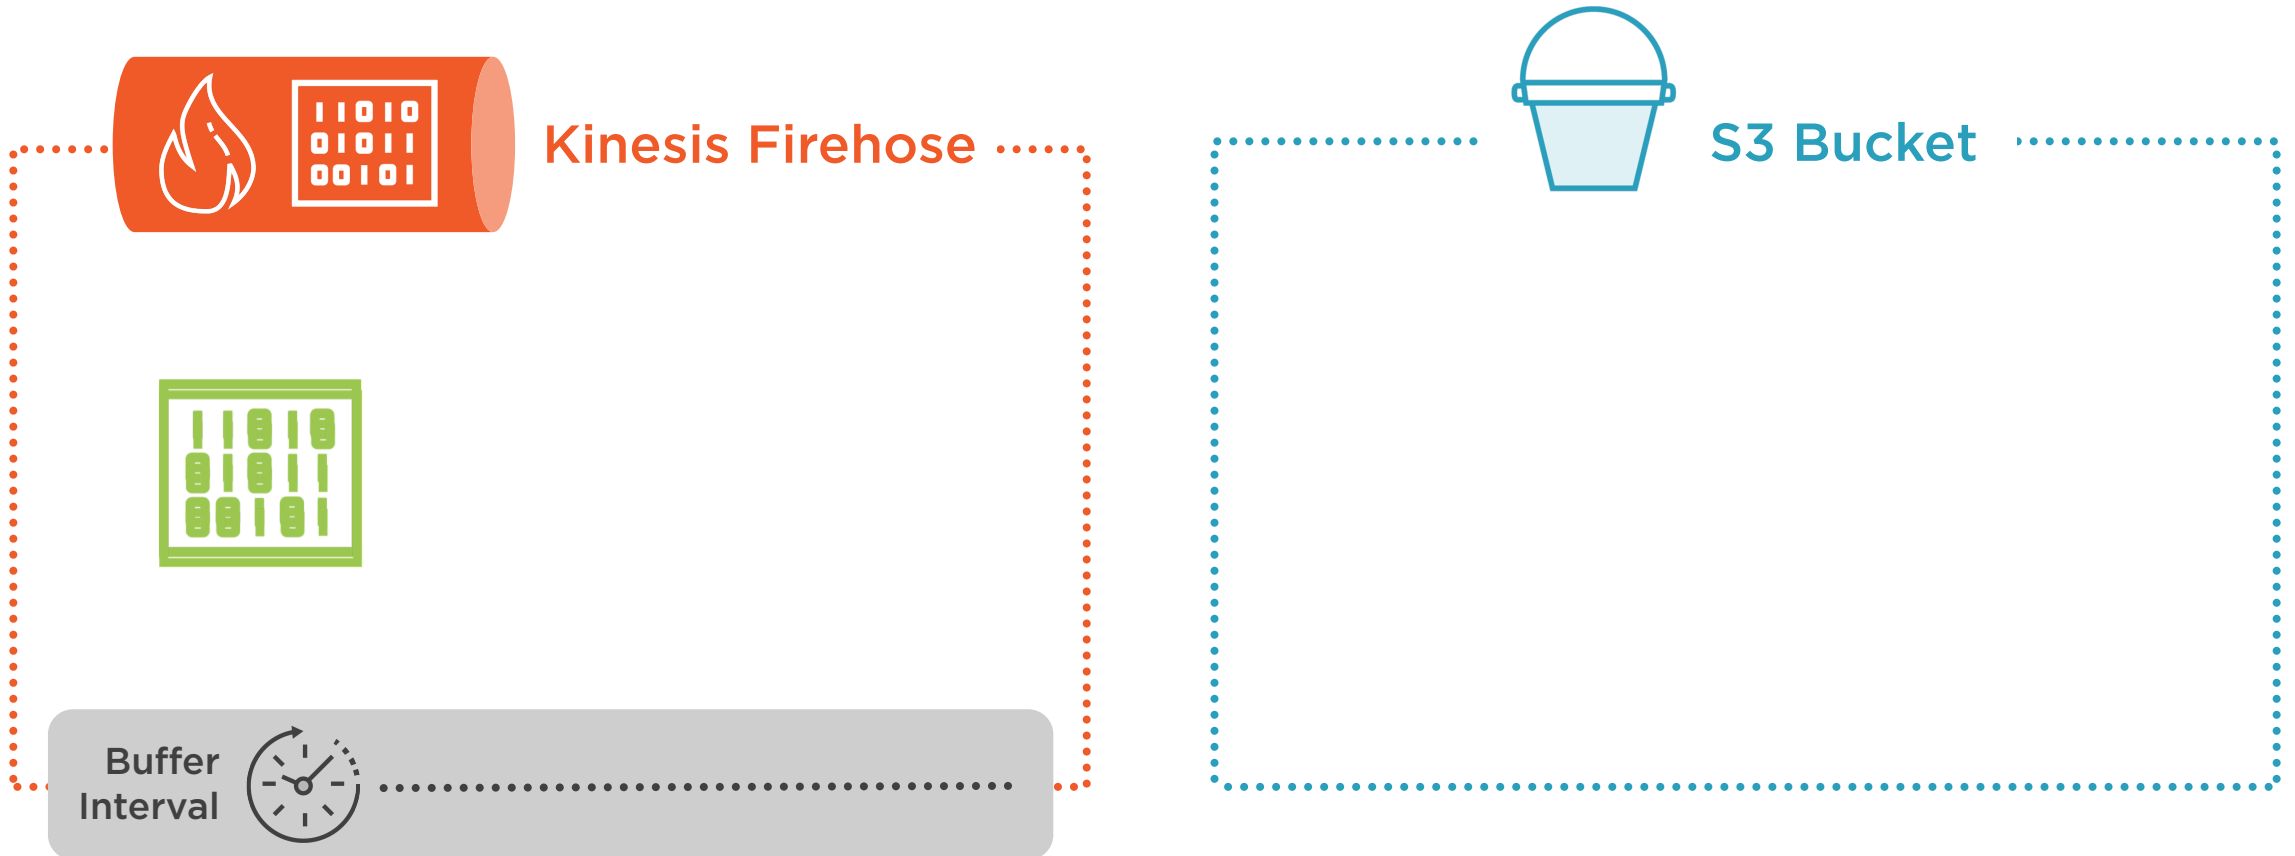
Buffer Size



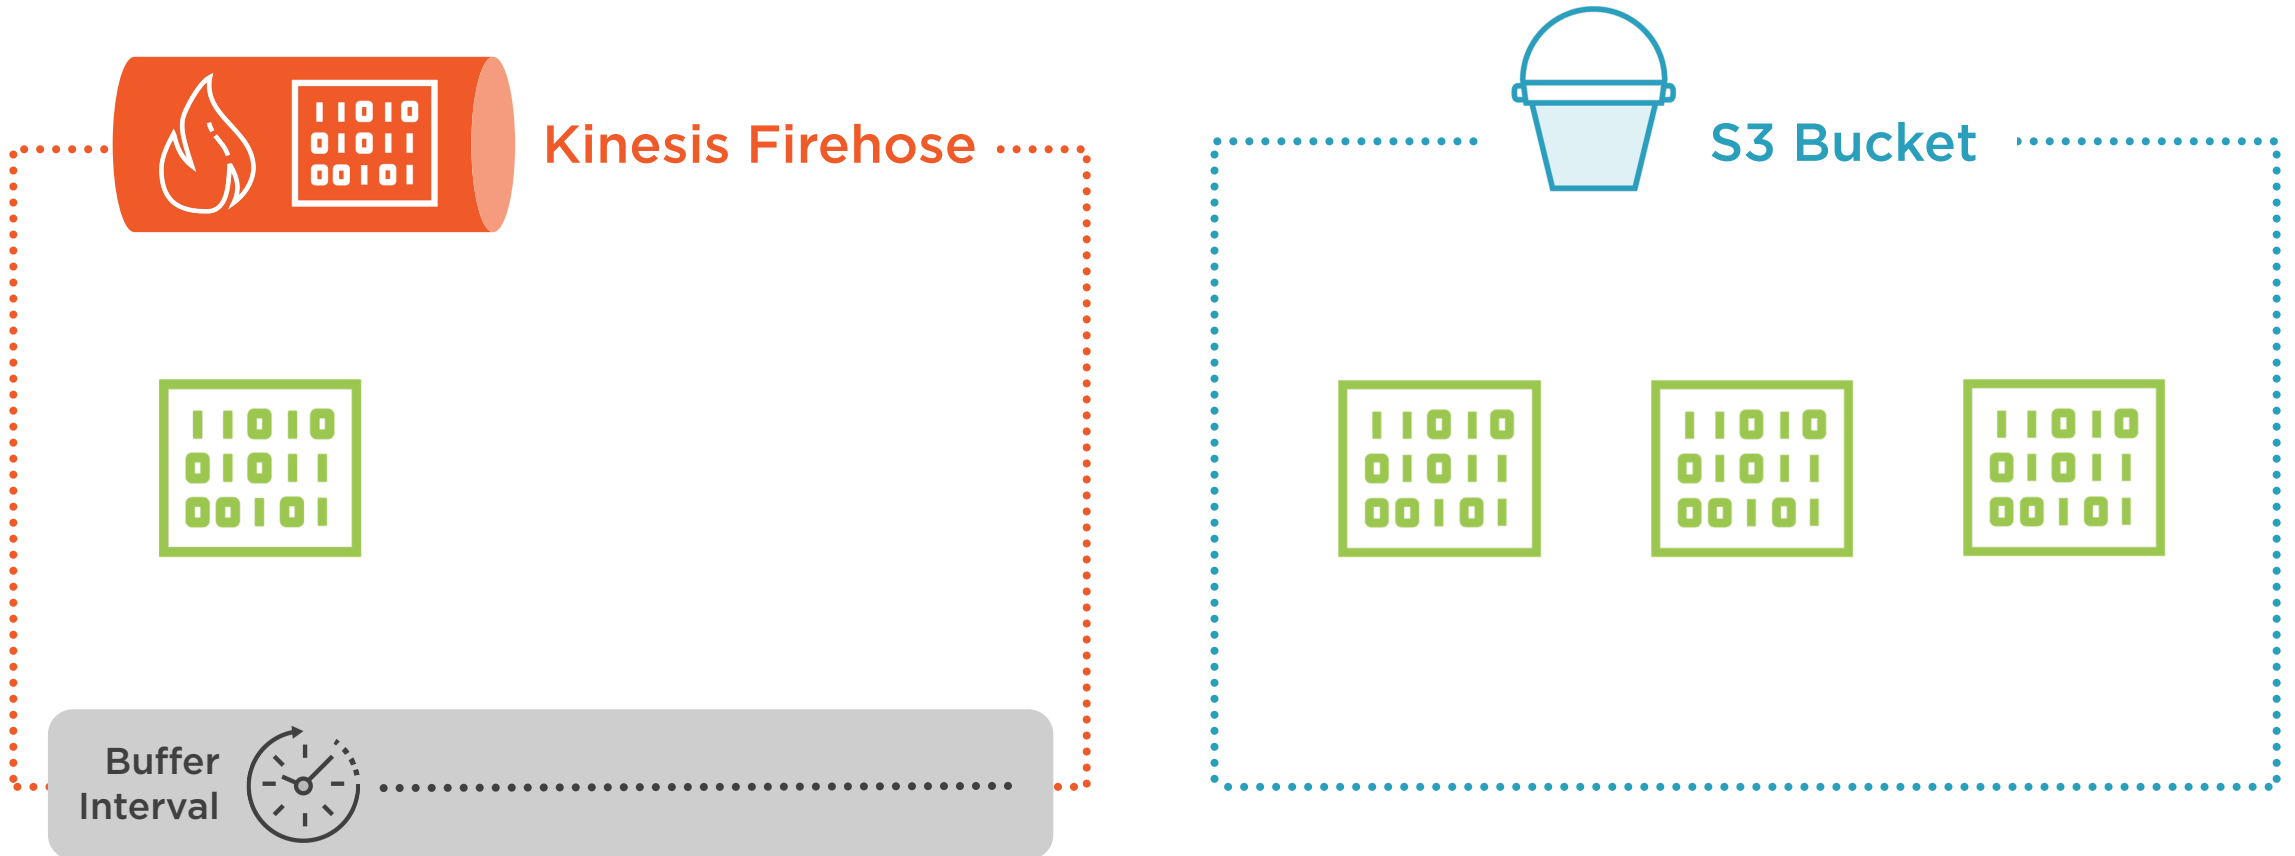
Buffer Size



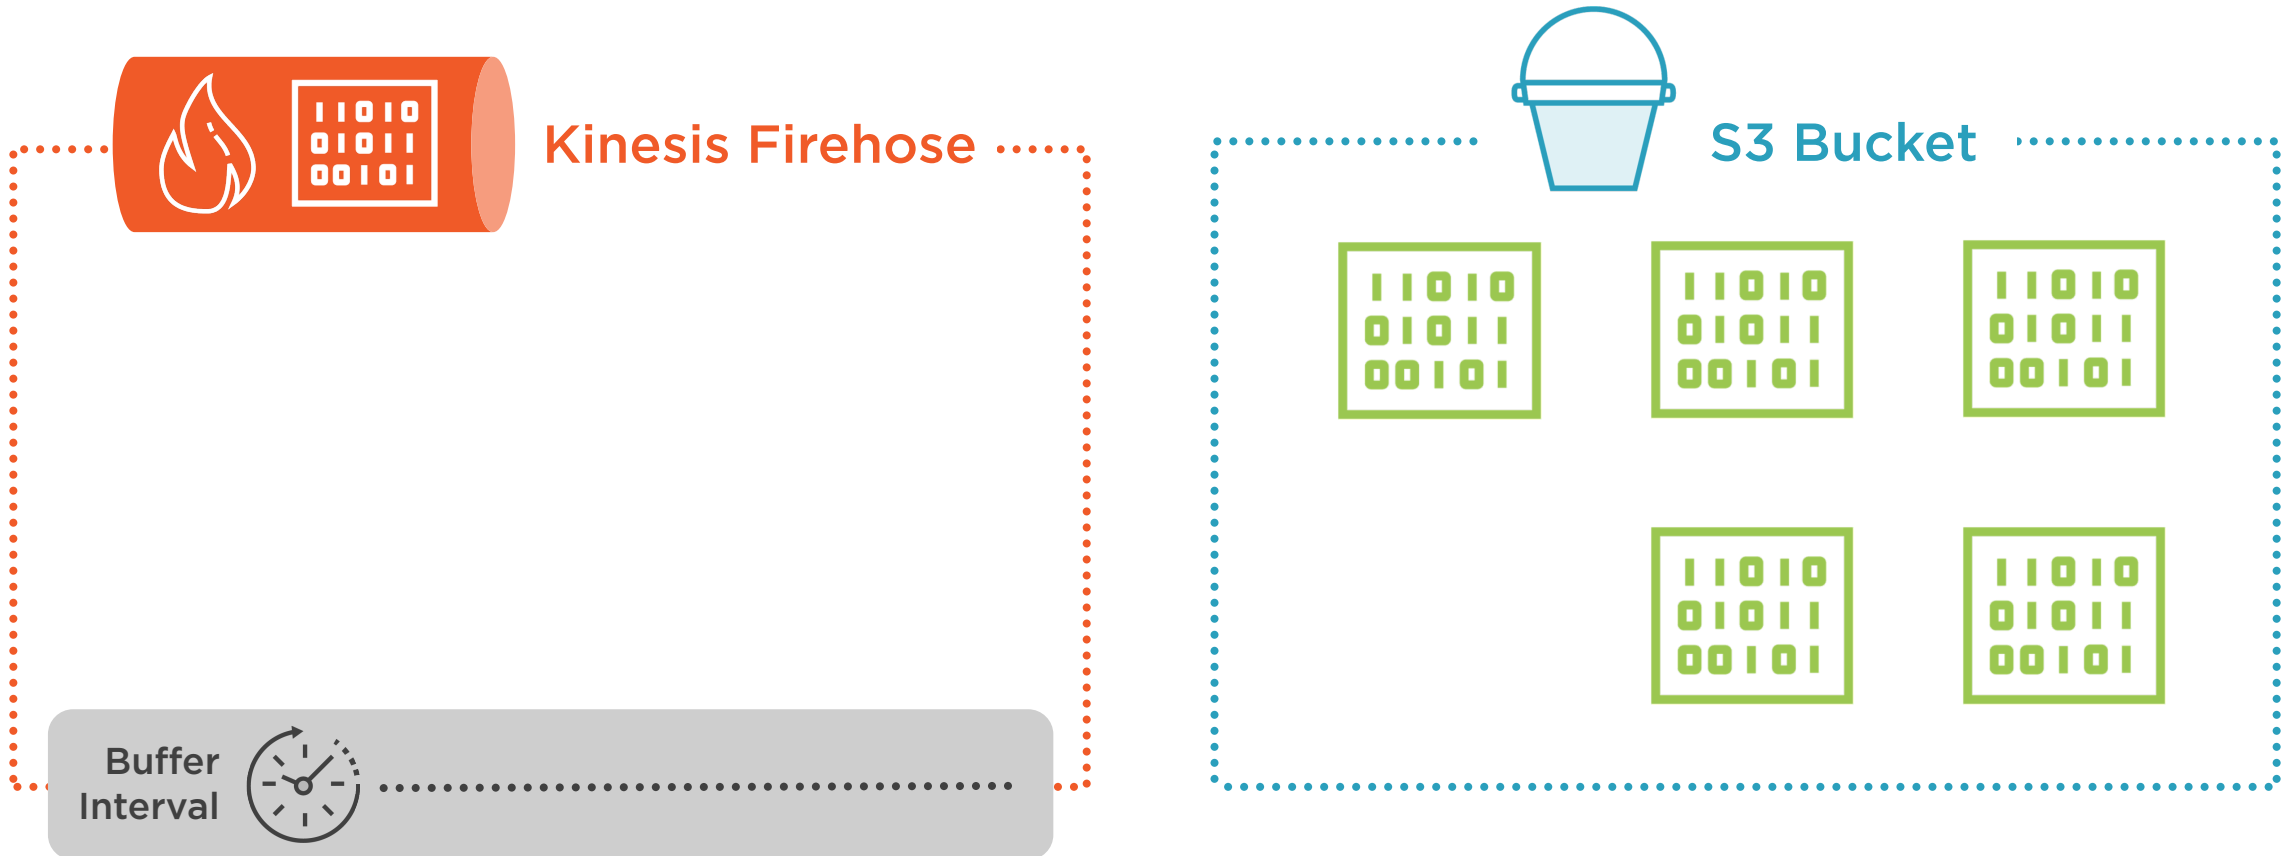
Buffer Interval



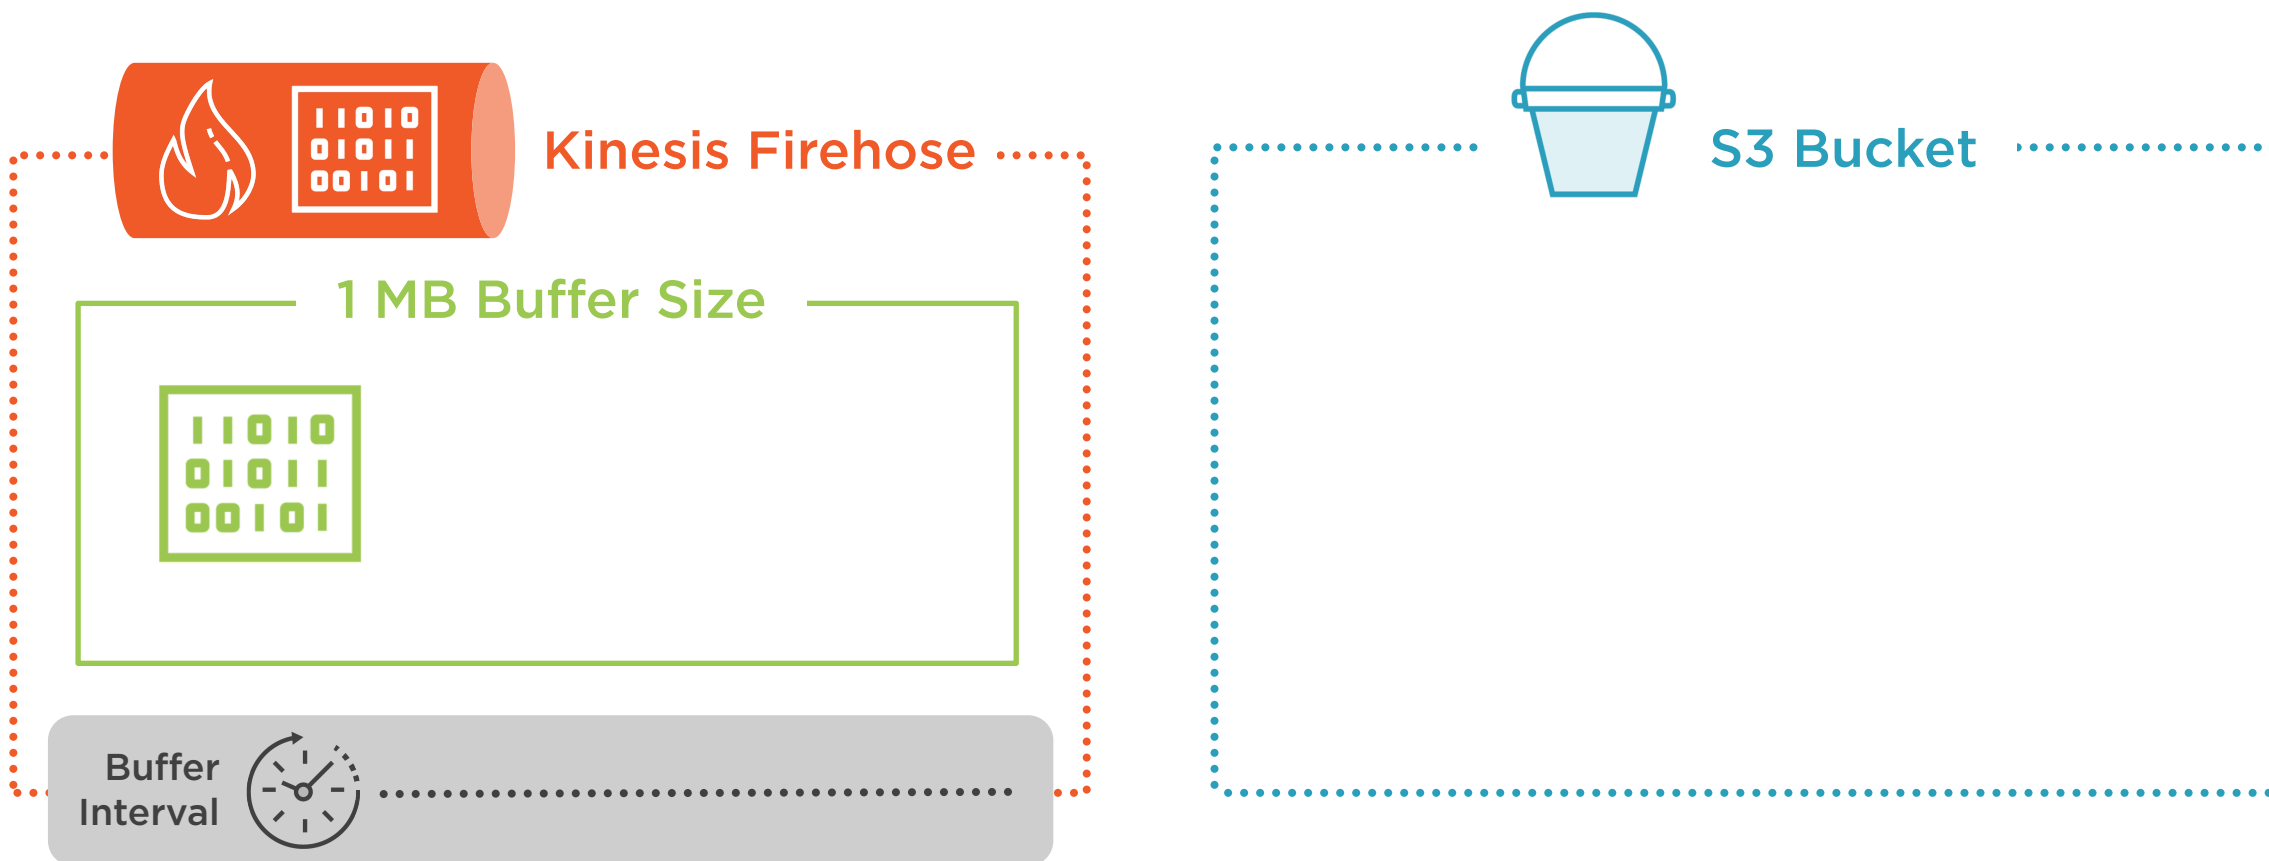
Buffer Interval



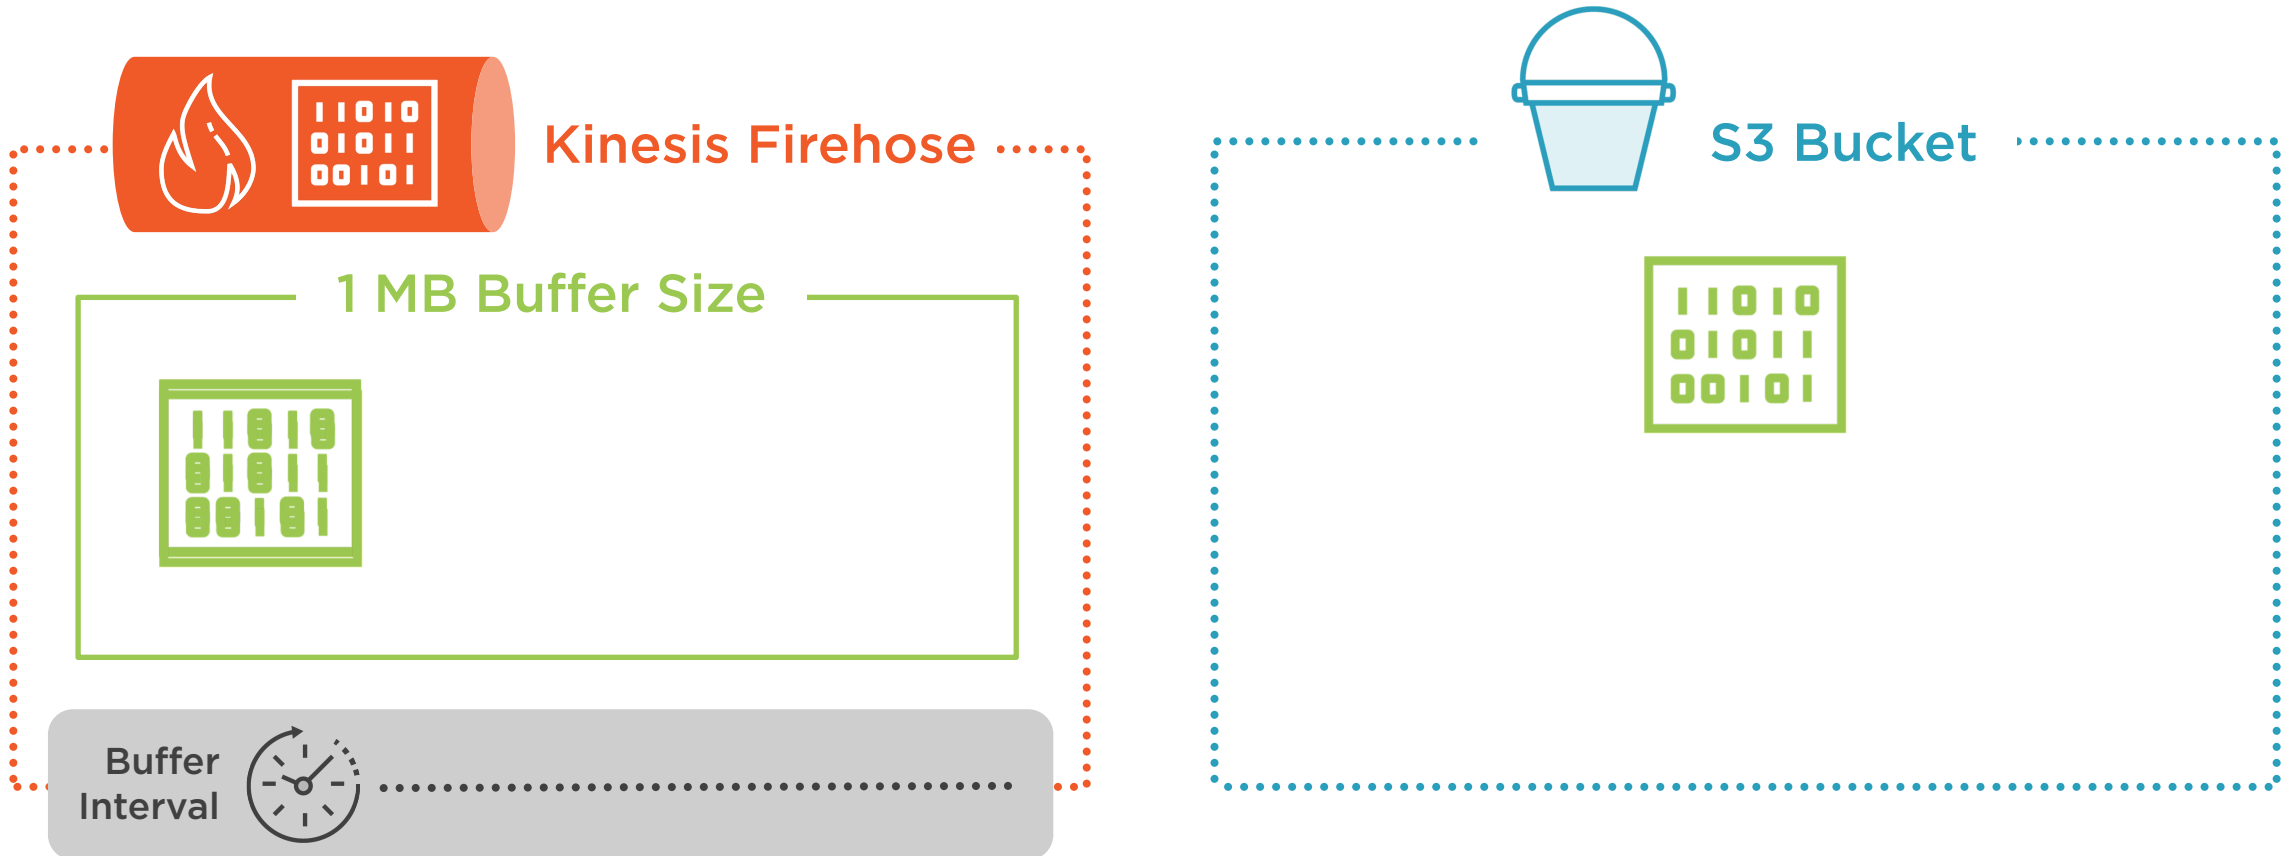
Buffer Interval



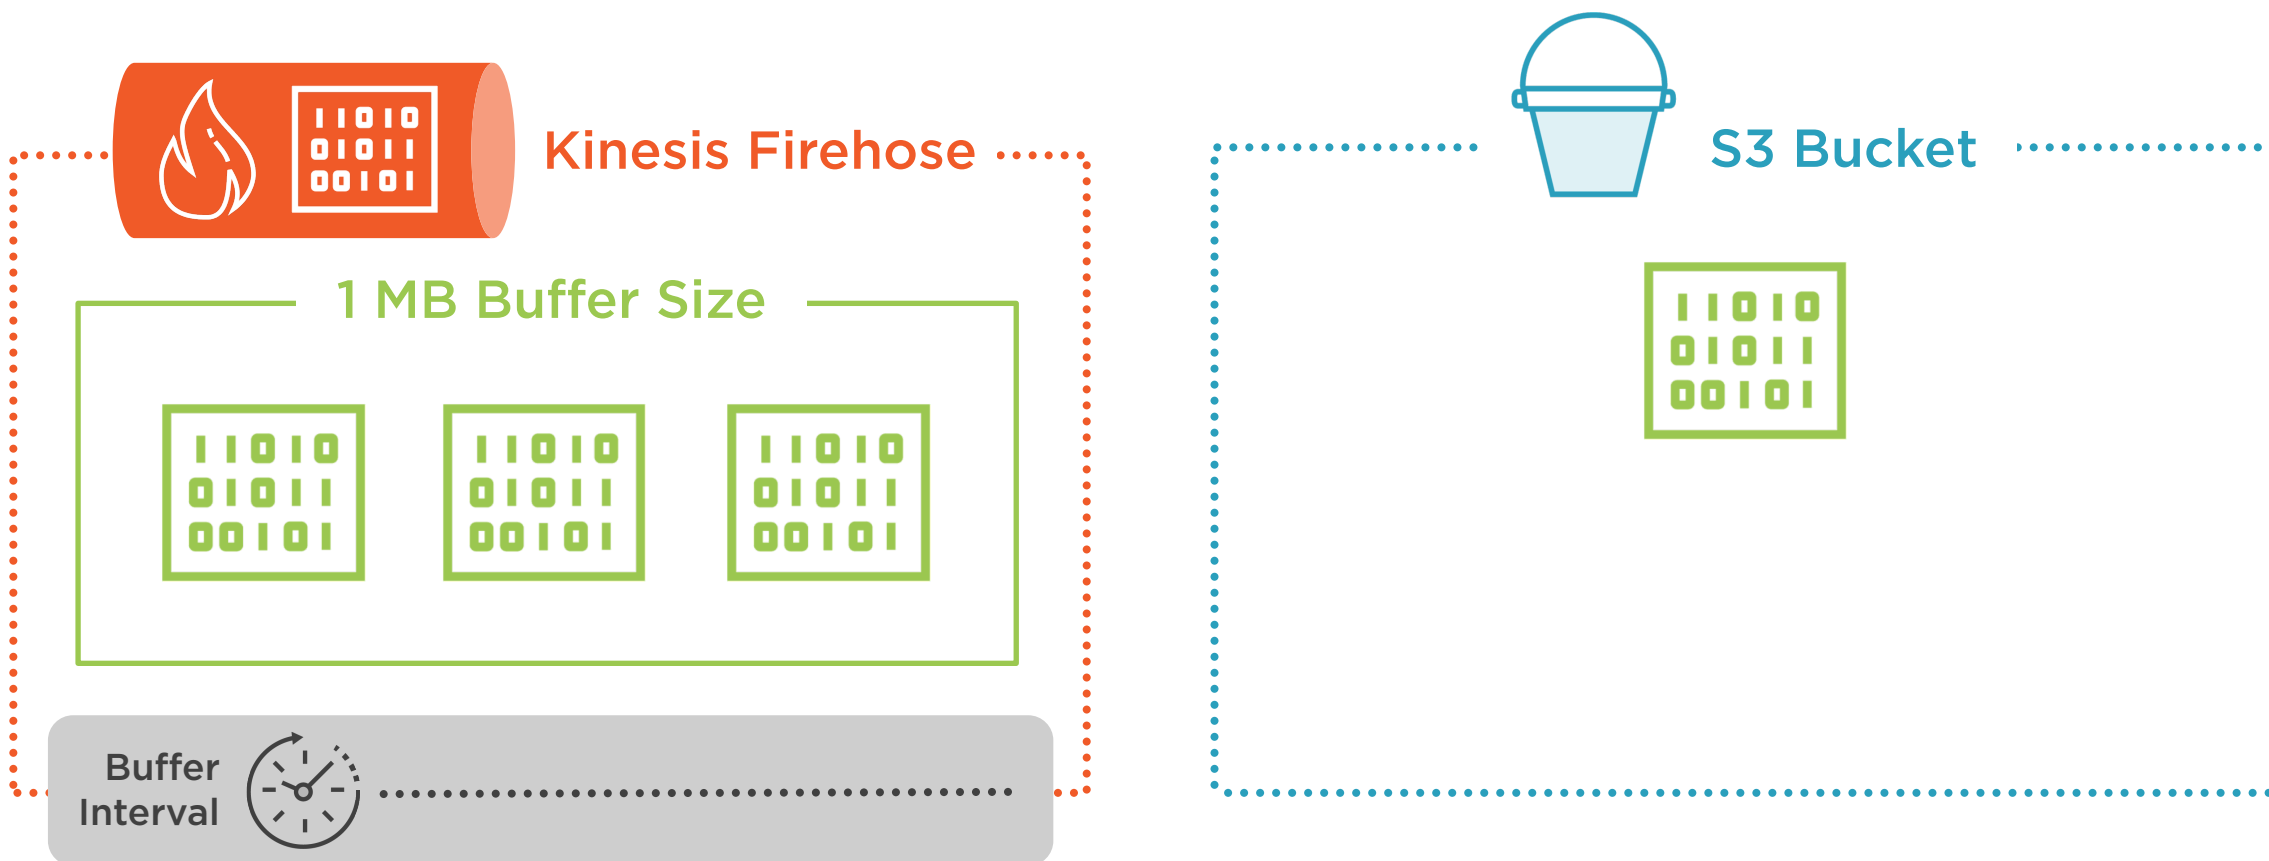
Buffer Size and Interval



Buffer Size and Interval



Buffer Size and Interval



Kinesis Data Firehose Limits



1k to 5k records/second (by region) – Soft limit



1k to 2k transactions/second (by region) – Soft limit



1MiB to 5MiB data ingested per second (by region) – Soft Limit



1000 KiB max record size – Hard limit



Kinesis Data Stream Producers – No ingestion limit




Kinesis Data Firehose: Transformations



Transformation Limitations



AWS Lambda only



Buffer size - 3MB

Transformation States



Ok: Processed correctly



Dropped: Intentionally dropped



Failed: Can be send to S3 for processing later



Let's visualize
some transformations



Data Source



Lambda



Kinesis Firehose

Buffer Size



Buffer
Interval



Data Source

Kinesis Firehose

S3 Bucket for Failures

Lambda



Data OK



Data Dropped



Data Failed



Buffer Size

Buffer
Interval




Amazon ES Destination




Kinesis Data Firehose: Producers



Kinesis Data Streams and Firehose Producers



Kinesis Agent



AWS SDKs




Kinesis Data Streams

Kinesis Firehose Producers



CloudWatch Log
Subscriptions



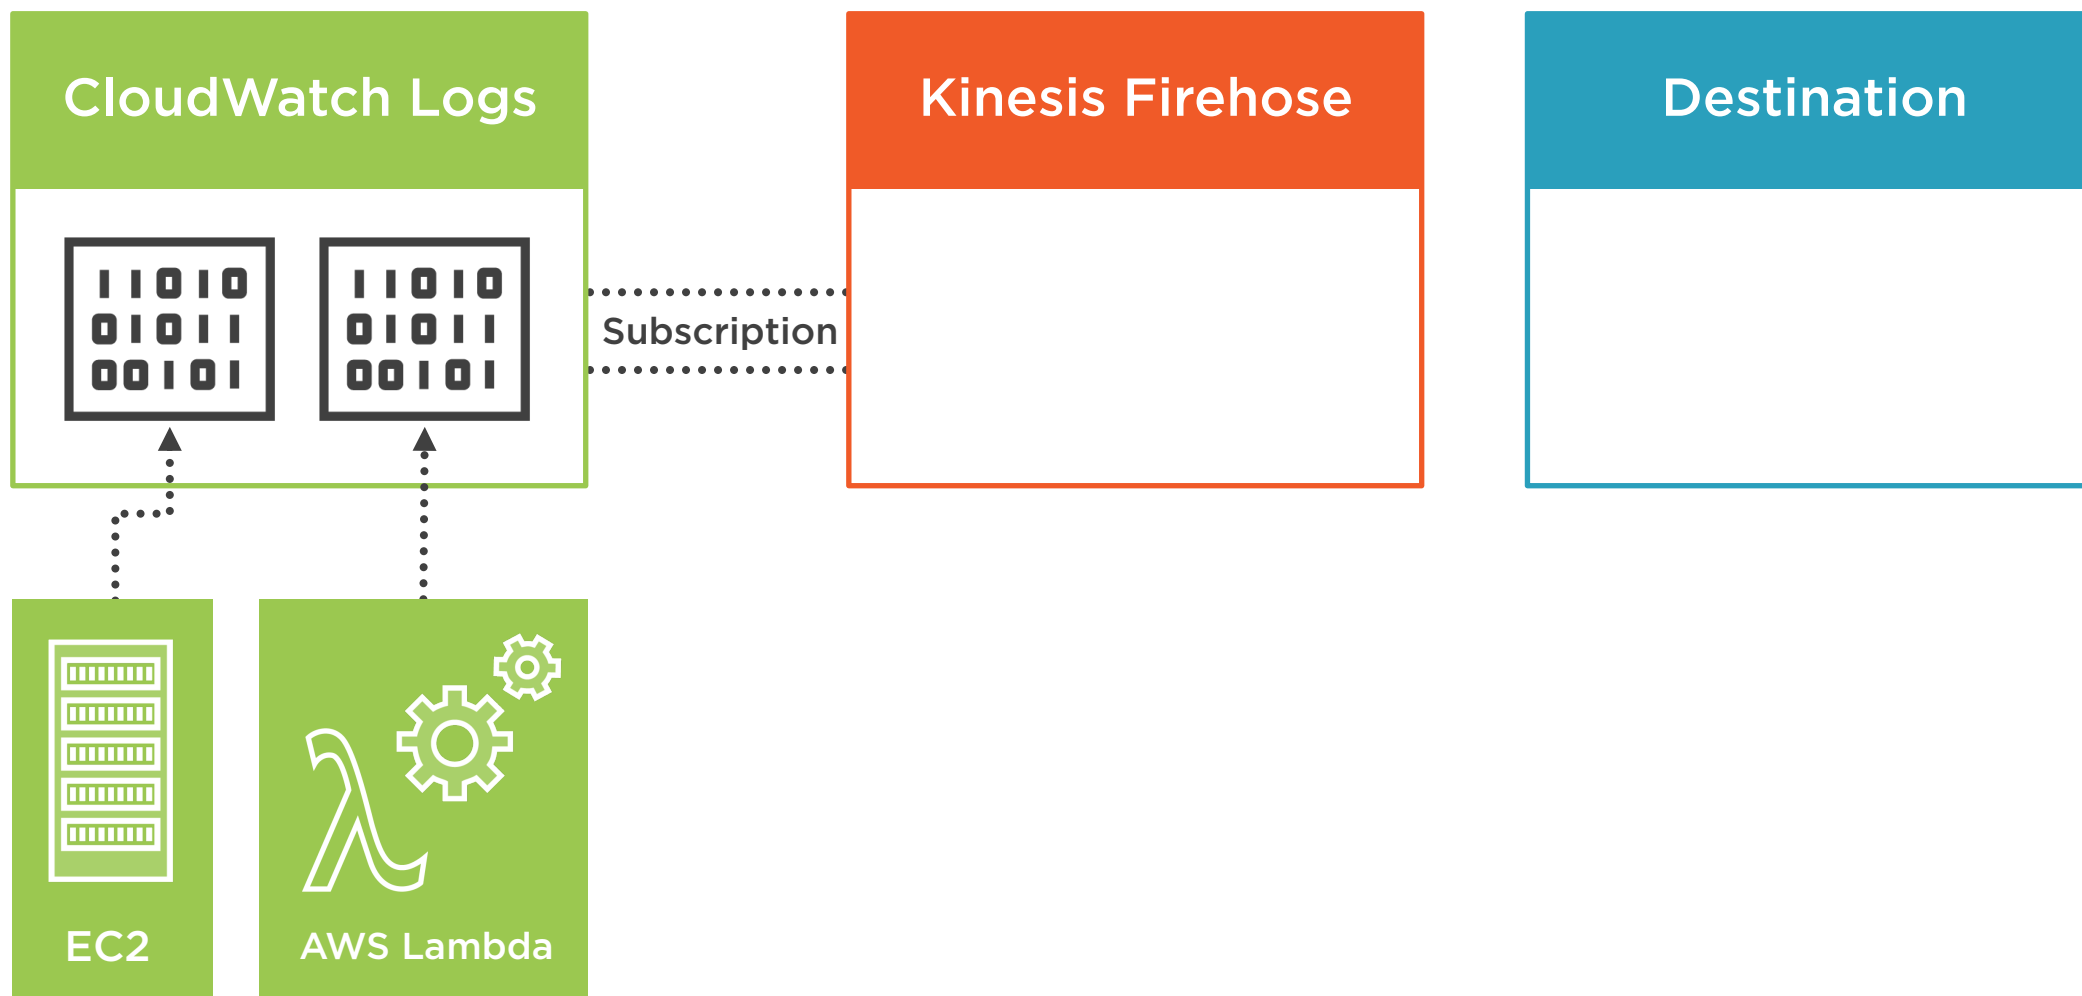
CloudWatch Logs
Events Rule



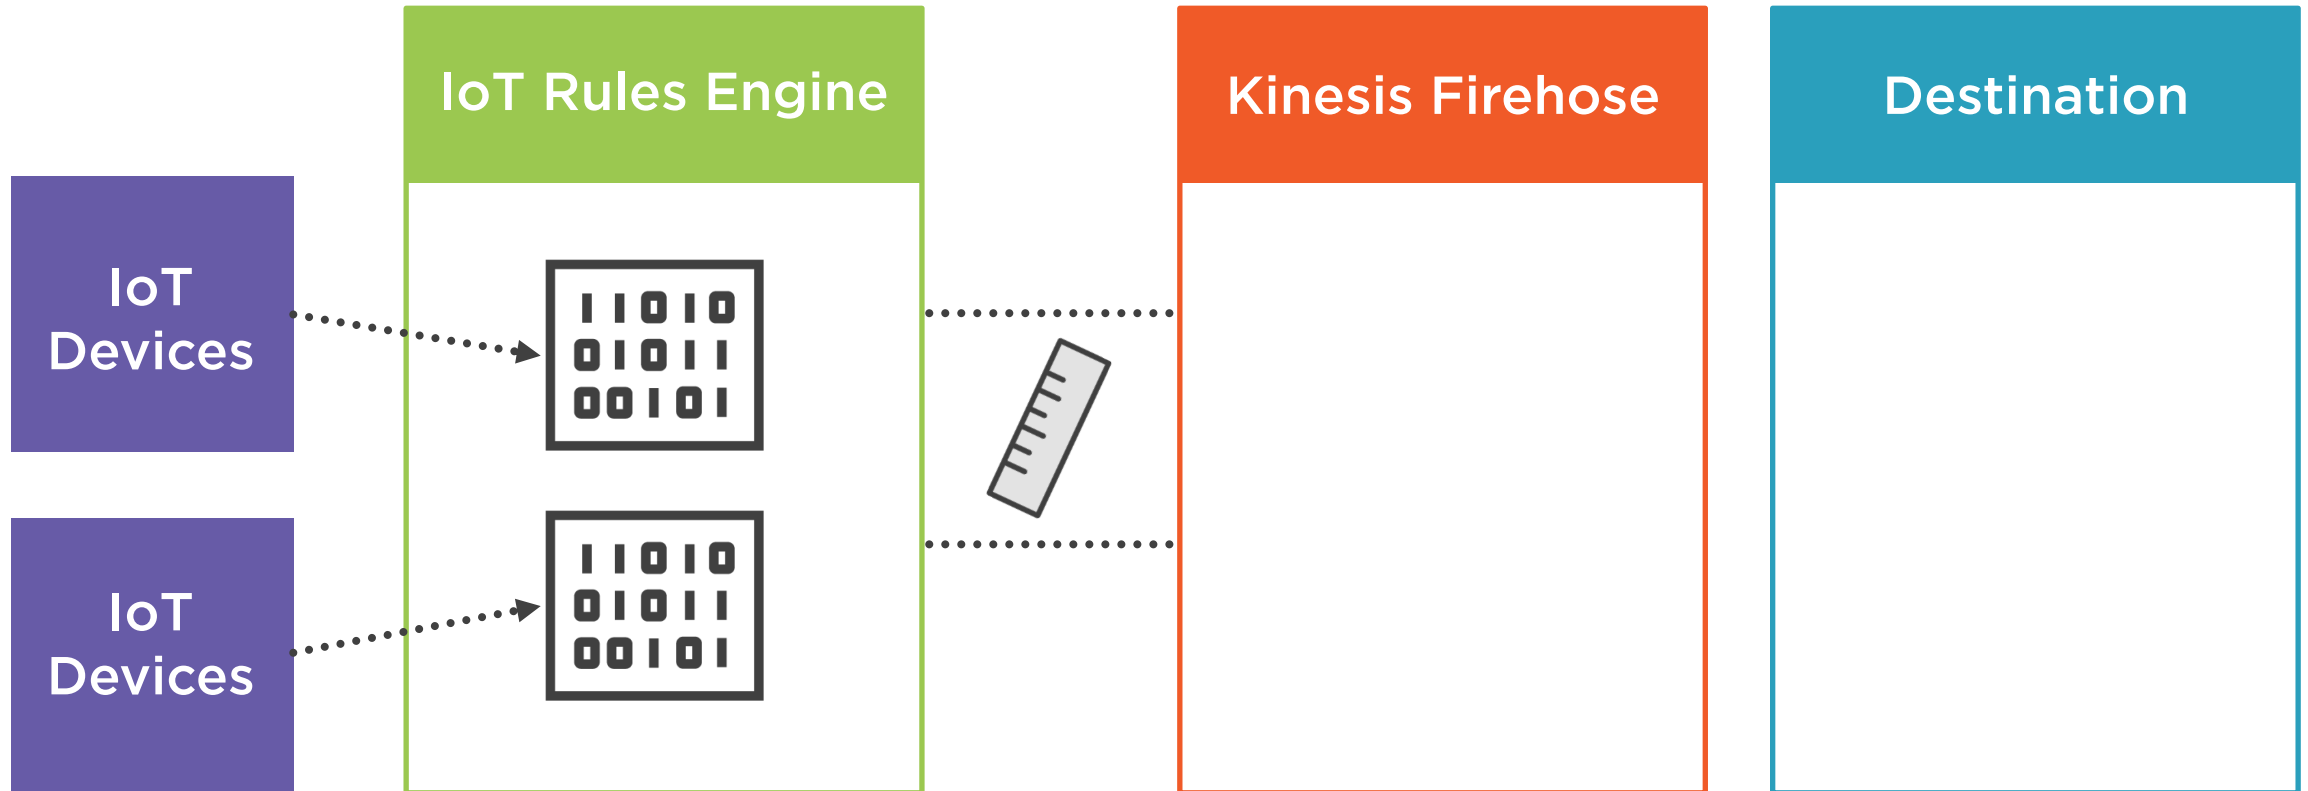
AWS IoT
Rules



CloudWatch Log Subscriptions



AWS IoT Rules



Kinesis Data Firehose: Destinations



Firehose Destinations



Amazon S3



Amazon ES



Redshift



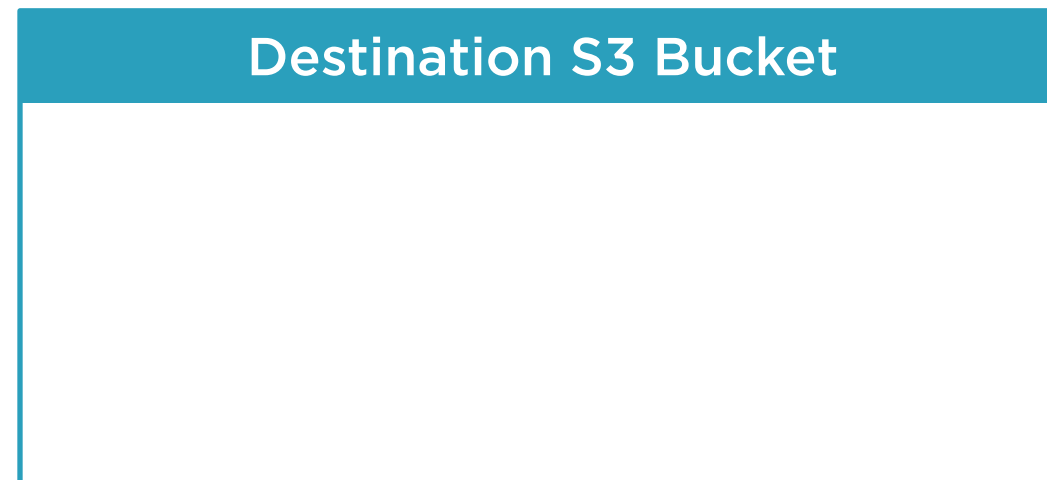
Splunk



Kinesis Firehose to Amazon S3



Kinesis Firehose to Amazon S3




Kinesis Firehose to Amazon S3

Kinesis Firehose

Buffer Interval 

Destination S3 Bucket

Lambda



Kinesis Firehose to Amazon S3

Kinesis Firehose



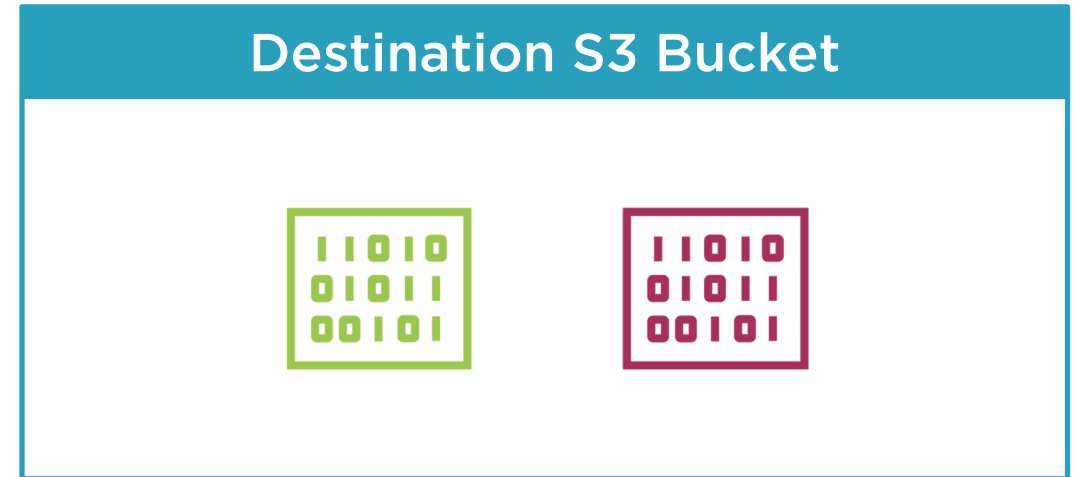
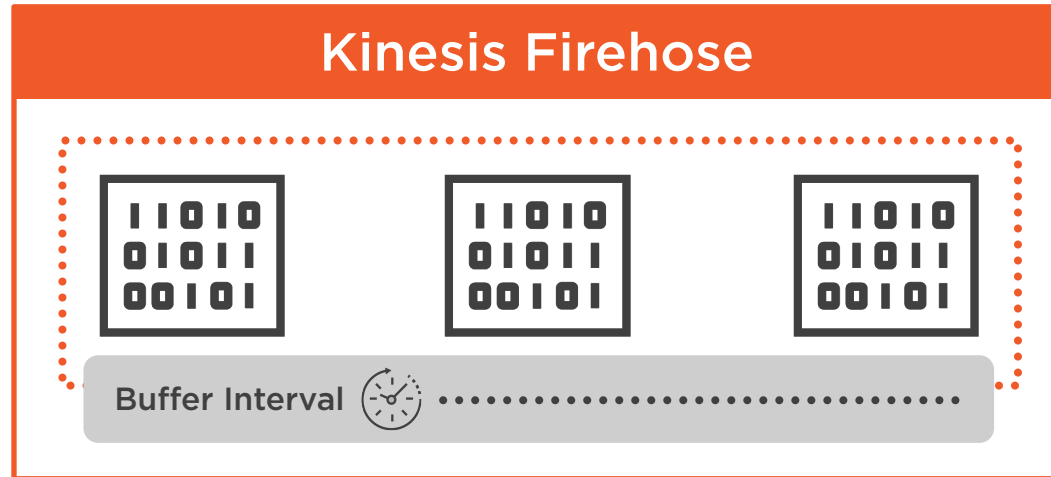


Destination S3 Bucket

Lambda



Kinesis Firehose to Amazon S3





Kinesis Firehose to Amazon S3

Kinesis Firehose




Buffer Interval 

Destination S3 Bucket

Lambda



Optional Backup Bucket



Kinesis Firehose to Amazon S3

Kinesis Firehose




Destination S3 Bucket

Lambda




Optional Backup Bucket



Kinesis Firehose to Amazon S3

Kinesis Firehose



Destination S3 Bucket



Lambda





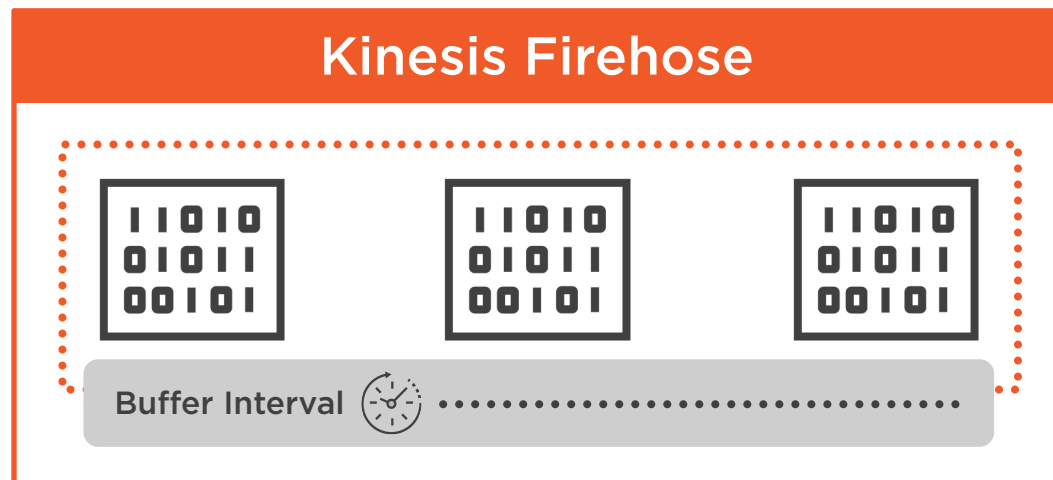
Optional Backup Bucket





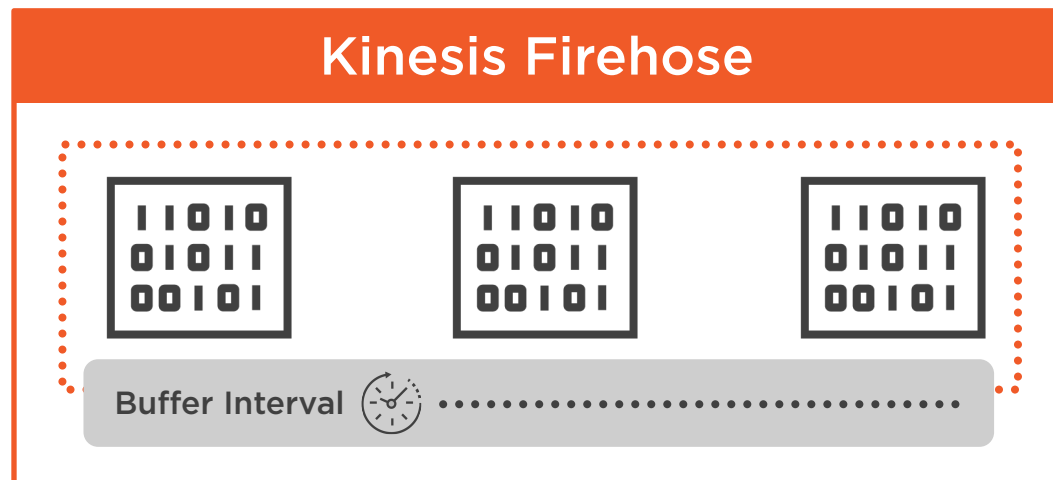
Kinesis Firehose to Amazon ES



Kinesis Firehose to Amazon ES




Kinesis Firehose to Amazon ES




Kinesis Firehose to Amazon ES

Kinesis Firehose







Amazon ES



S3 Bucket



Kinesis Firehose to Amazon ES





Kinesis Firehose to Amazon ES

Kinesis Firehose

Buffer Interval 

Amazon ES

Lambda



S3 Bucket



Kinesis Firehose to Amazon ES

Kinesis Firehose



Amazon ES

Lambda



S3 Bucket



Kinesis Firehose to Amazon ES


Kinesis Firehose

Amazon ES




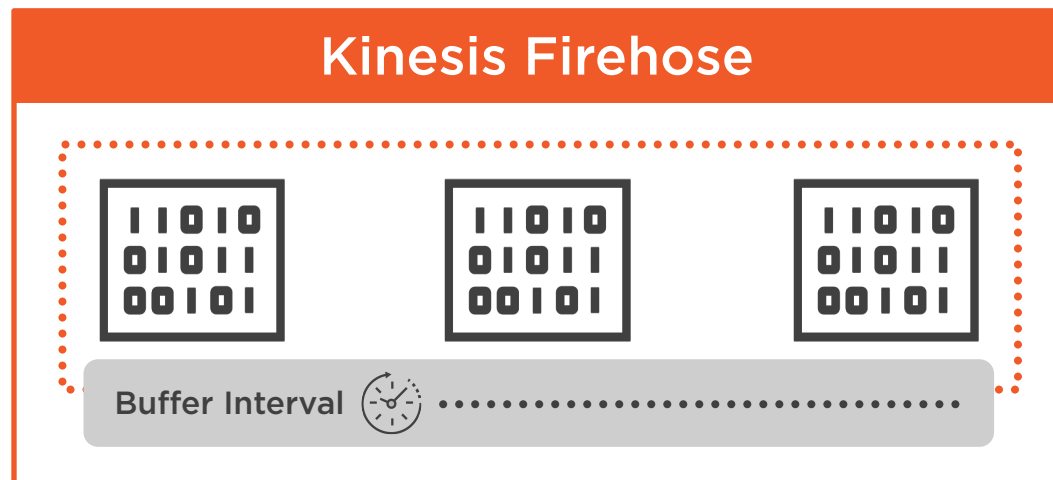


Lambda

S3 Bucket



Kinesis Firehose to Amazon ES



Kinesis Firehose to Amazon ES

Kinesis Firehose

11010
01011
00101

11010
01011
00101

11010
01011
00101

Buffer Interval 

Amazon ES

Lambda

S3 Bucket

11010
01011
00101

11010
01011
00101

11010
01011
00101






Kinesis Firehose to Amazon ES

Kinesis Firehose




Buffer Interval 

Amazon ES



Lambda




S3 Bucket





Kinesis Firehose to Amazon Redshift



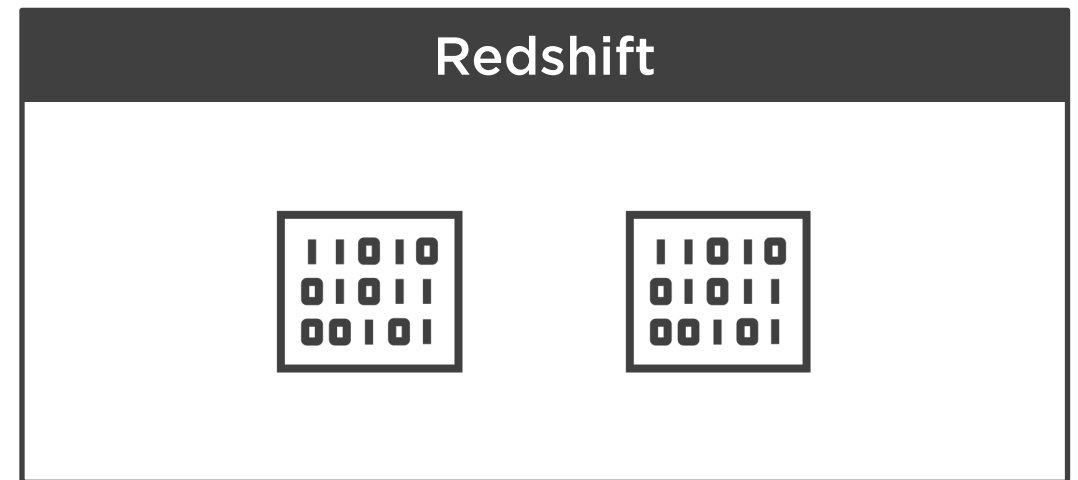


Kinesis Firehose to Amazon Redshift



Kinesis Firehose to Amazon Redshift



Kinesis Firehose to Amazon Redshift





Kinesis Firehose to Amazon Redshift

Kinesis Firehose





Intermediate Bucket



Redshift



Kinesis Firehose to Amazon Redshift





Kinesis Firehose to Amazon Redshift

Kinesis Firehose

Buffer Interval 

Intermediate Bucket

Lambda




Redshift



Kinesis Firehose to Amazon Redshift

Kinesis Firehose



Intermediate Bucket

Lambda




Redshift



Kinesis Firehose to Amazon Redshift

Kinesis Firehose

Intermediate Bucket

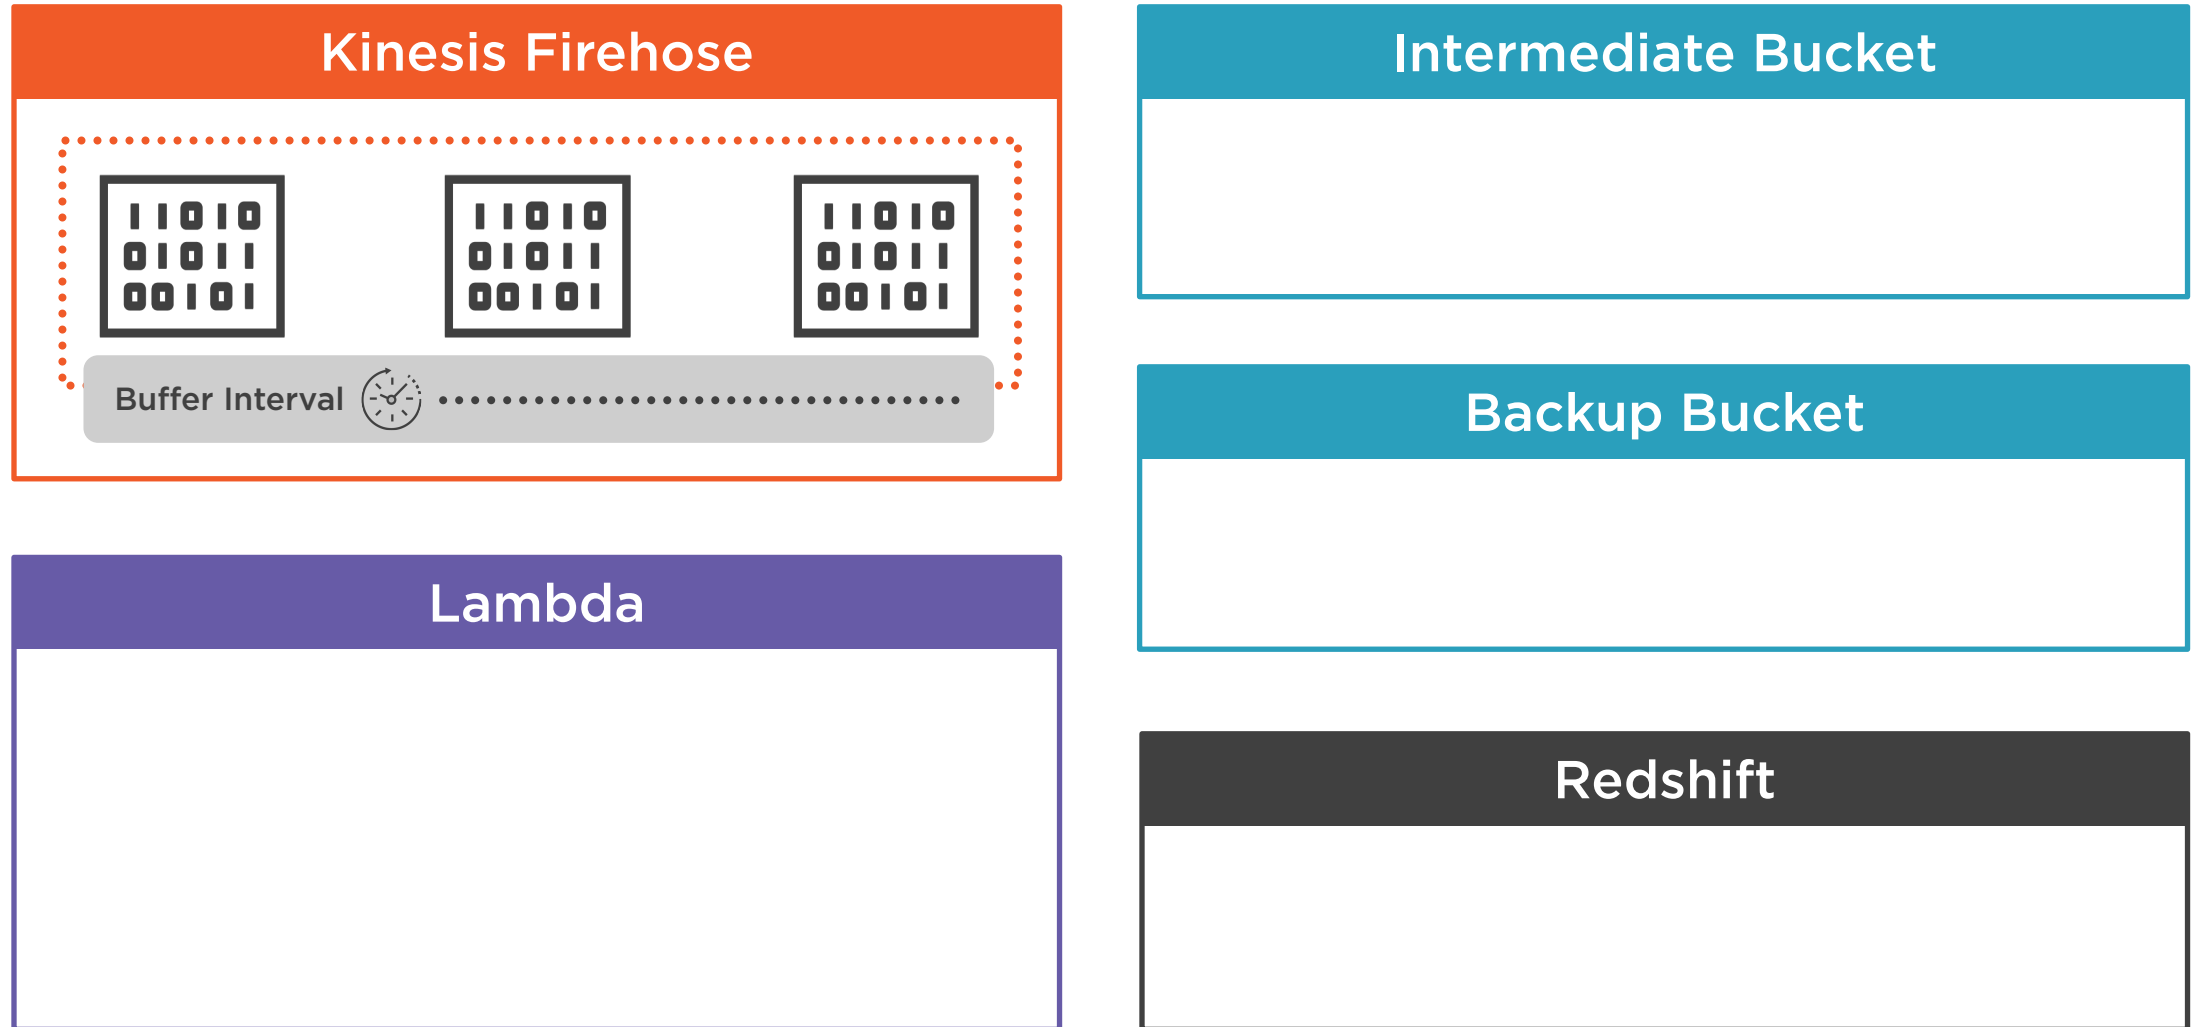


Lambda

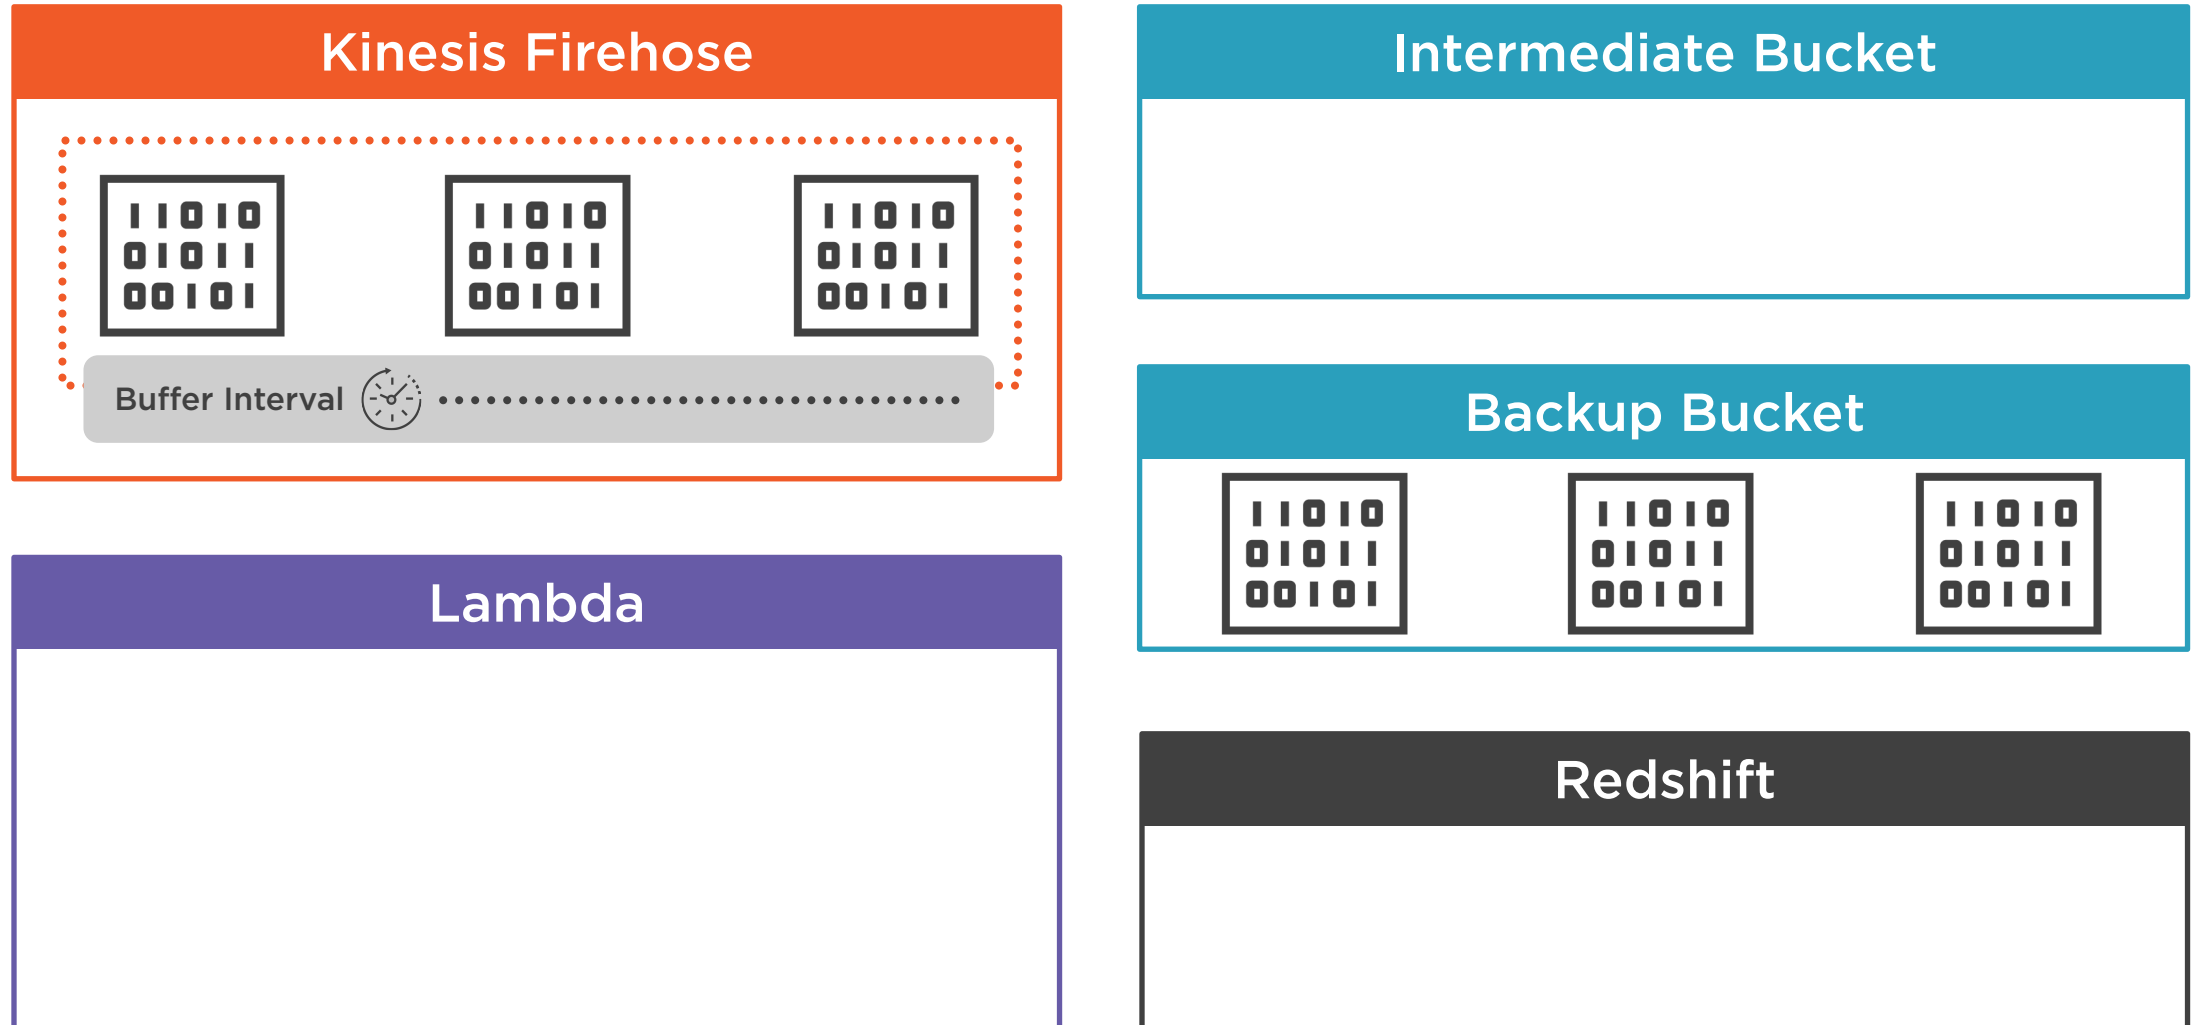
Redshift







Kinesis Firehose to Amazon Redshift








Kinesis Firehose to Amazon Redshift






Kinesis Firehose to Amazon Redshift



Kinesis Firehose to Amazon Redshift





Kinesis Firehose to Amazon Redshift





Kinesis Firehose to Splunk






Kinesis Firehose to Splunk



Kinesis Firehose to Splunk



Kinesis Firehose to Splunk





Kinesis Firehose to Splunk

Kinesis Firehose

Buffer Interval 

Splunk

Lambda



S3 Bucket





Kinesis Firehose to Splunk

Kinesis Firehose

Splunk


Lambda

S3 Bucket



Kinesis Firehose to Splunk

Kinesis Firehose



Splunk

Lambda



S3 Bucket



Kinesis Firehose to Splunk

Kinesis Firehose

Splunk






Lambda




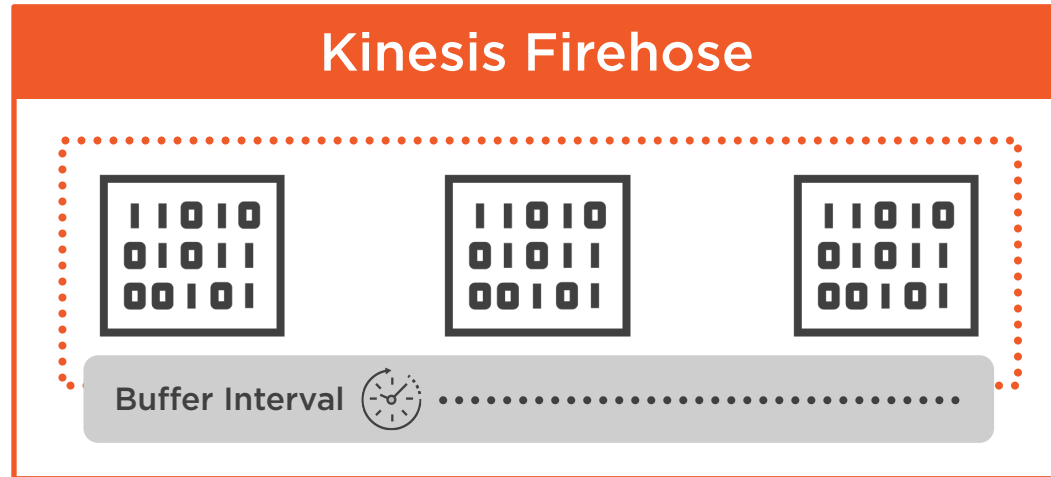
S3 Bucket



Kinesis Firehose to Splunk



Kinesis Firehose to Splunk





Kinesis Firehose to Splunk

Kinesis Firehose



Buffer Interval 

Splunk

Lambda




S3 Bucket



Kinesis Firehose to Splunk

Kinesis Firehose




Splunk

Lambda



S3 Bucket



Kinesis Firehose to Splunk




Kinesis Firehose

Splunk



Lambda

S3 Bucket



Demo



WARNING

- Not in the AWS Free Tier

Creating a Firehose Delivery Stream

- Create Kinesis Data Stream Source
- Script to send data into data stream
- Create Firehose Delivery Stream
- Create Lambda function
 - Transform data from JSON to CSV
- Output data in S3



Kinesis Firehose vs. Kinesis Data Streams



Comparing Kinesis Options

Kinesis Firehose

1-minute delay for data is ok

More managed service

Only supports specific destinations

Use Lambda for data transformation

Retries deliveries automatically

Kinesis Data Streams

You want data in near real time

More control of capacity yourself

Supports multiple data consumers

Consumers processing data streams

24 hours to 7 days of data retention



Summary



Basic Concepts

Firehose Producers and Destinations

Transforming data

Demo: Working with Firehose

Firehose vs. Data Streams





Up next:

Data Transfer with Snowball and Direct Connect

